

Install Packages

```
In [1]: 1 #pip install opencv-python
```

```
In [2]: 1 #pip install -U seaborn
```

```
In [3]: 1 !pip install shap
```

```
Requirement already satisfied: shap in c:\users\peddi\anaconda3\lib\site-packages (0.39.0)
Requirement already satisfied: scipy in c:\users\peddi\anaconda3\lib\site-packages (from shap) (1.5.2)
Requirement already satisfied: cloudpickle in c:\users\peddi\anaconda3\lib\site-packages (from shap) (1.6.0)
Requirement already satisfied: slicer==0.0.7 in c:\users\peddi\anaconda3\lib\site-packages (from shap) (0.0.7)
Requirement already satisfied: numba in c:\users\peddi\anaconda3\lib\site-packages (from shap) (0.51.2)
Requirement already satisfied: scikit-learn in c:\users\peddi\anaconda3\lib\site-packages (from shap) (0.22.2)
Requirement already satisfied: pandas in c:\users\peddi\anaconda3\lib\site-packages (from shap) (1.1.3)
Requirement already satisfied: tqdm>4.25.0 in c:\users\peddi\anaconda3\lib\site-packages (from shap) (4.50.2)
Requirement already satisfied: numpy in c:\users\peddi\anaconda3\lib\site-packages (from shap) (1.19.2)
Requirement already satisfied: llvmlite<0.35,>=0.34.0.dev0 in c:\users\peddi\anaconda3\lib\site-packages (from numba->shap) (0.34.0)
Requirement already satisfied: setuptools in c:\users\peddi\anaconda3\lib\site-packages (from numba->shap) (50.3.1.post20201107)
Requirement already satisfied: joblib>=0.11 in c:\users\peddi\anaconda3\lib\site-packages (from scikit-learn->shap) (0.17.0)
Requirement already satisfied: pytz>=2017.2 in c:\users\peddi\anaconda3\lib\site-packages (from pandas->shap) (2020.1)
Requirement already satisfied: python-dateutil>=2.7.3 in c:\users\peddi\anaconda3\lib\site-packages (from pandas->shap) (2.8.1)
Requirement already satisfied: six>=1.5 in c:\users\peddi\anaconda3\lib\site-packages (from python-dateutil>=2.7.3->pandas->shap) (1.15.0)
```

```
In [4]: 1 import shap
```

In [5]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import cv2
6 import tensorflow as tf
7 from PIL import Image
8 import os
9 from sklearn.model_selection import train_test_split
10 from keras.utils import to_categorical
11 from keras.models import Sequential, load_model
12 from keras.layers import Conv2D, MaxPool2D, Dense, Flatten, Dropout
13 import matplotlib.image as mpimg
14 from matplotlib import *
```

In [165]:

```
1 from tensorflow import keras
2 from tensorflow.keras import layers
3 from tensorflow.keras.models import Sequential
4 from keras.models import Sequential
5 from tensorflow.keras.layers import BatchNormalization
6 from keras.layers import Input, Flatten, Dense, Dropout, Conv2D, concatenate
7 from tensorflow.keras.activations import relu, softmax
8 from tensorflow.keras.initializers import he_normal, zeros, glorot_normal, RandomNormal
9 from tensorflow.keras.models import Model
10 from tensorflow.keras.optimizers import Adam
11 from tensorflow.keras.losses import SparseCategoricalCrossentropy
12 from tensorflow.keras.regularizers import l1, l2, l1_l2
```

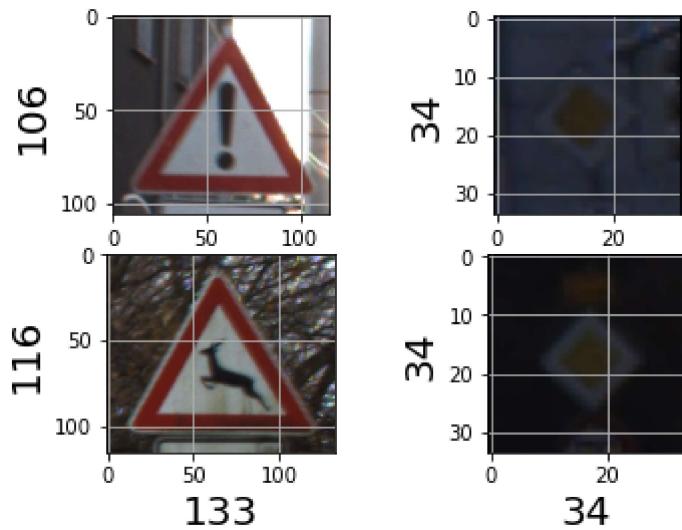
In [7]:

```
1 # Metrics
2 from sklearn.metrics import classification_report
3 from sklearn.metrics import confusion_matrix
4 from sklearn.metrics import accuracy_score
```

Extract Data from Images

In [28]:

```
1 import random
2 from matplotlib.image import imread
3 dataset_dir = 'C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Analytics\\\\Fin
4 meta_info = os.path.join(dataset_dir, 'Meta.csv')
5 train_csv_path = os.path.join(dataset_dir, 'Train.csv')
6 test_csv_path = os.path.join(dataset_dir, 'Test.csv')
7
8 test = pd.read_csv(dataset_dir + 'Train.csv')
9 imgs = test["Path"].values
10 imgs
11
12 # Visualizing 25 random images from train data
13 import random
14 from matplotlib.image import imread
15 dataset_dir = 'C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Analytics\\\\Fin
16 meta_info = os.path.join(dataset_dir, 'Meta.csv')
17 train_csv_path = os.path.join(dataset_dir, 'Train.csv')
18 test_csv_path = os.path.join(dataset_dir, 'Test.csv')
19
20 test = pd.read_csv(dataset_dir + 'Train.csv')
21 imgs = test["Path"].values
22
23 #plt.figure(figsize=(25,25))
24
25 for i in range(1,5):
26     plt.subplot(2,2,i)
27     random_img_path = dataset_dir + '/' + random.choice(imgs)
28     rand_img = imread(random_img_path)
29     plt.imshow(rand_img)
30     plt.grid(b=None)
31     plt.xlabel(rand_img.shape[1], fontsize = 20)#width of image
32     plt.ylabel(rand_img.shape[0], fontsize = 20)#height of image
```



In [29]:

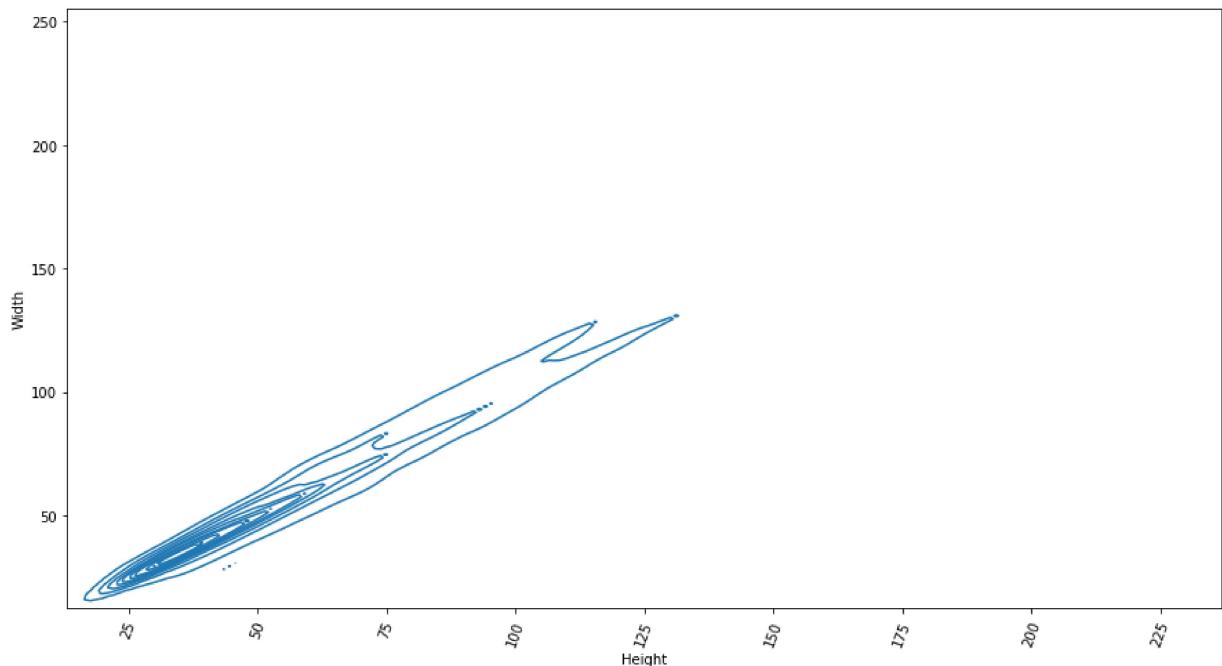
```
1 train_data_color = '#0f7b8e'
2 test_data_color = '#630f8e'
3
4 df_train = pd.read_csv(train_csv_path)
5 df_test = pd.read_csv(test_csv_path)
6 df_meta = pd.read_csv(meta_info)
7
8 df_train['Path'] = list(map(lambda x: os.path.join(dataset_dir,x.lower()), df_train['Path']))
9 df_test['Path'] = list(map(lambda x: os.path.join(dataset_dir,x.lower()), df_test['Path']))
10 df_meta['Path'] = list(map(lambda x: os.path.join(dataset_dir,x.lower()), df_meta['Path']))
11
12 df_train.sample(3)
```

Out[29]:

	Width	Height	Roi.X1	Roi.Y1	Roi.X2	Roi.Y2	ClassId	Path
36574	32	32	6	6	27	27	38	C:\Drive\Siri\MSDSBA\4. IE7860 Intelligent Ana...
34811	39	39	6	6	34	34	35	C:\Drive\Siri\MSDSBA\4. IE7860 Intelligent Ana...
35543	46	48	6	6	41	43	36	C:\Drive\Siri\MSDSBA\4. IE7860 Intelligent Ana...

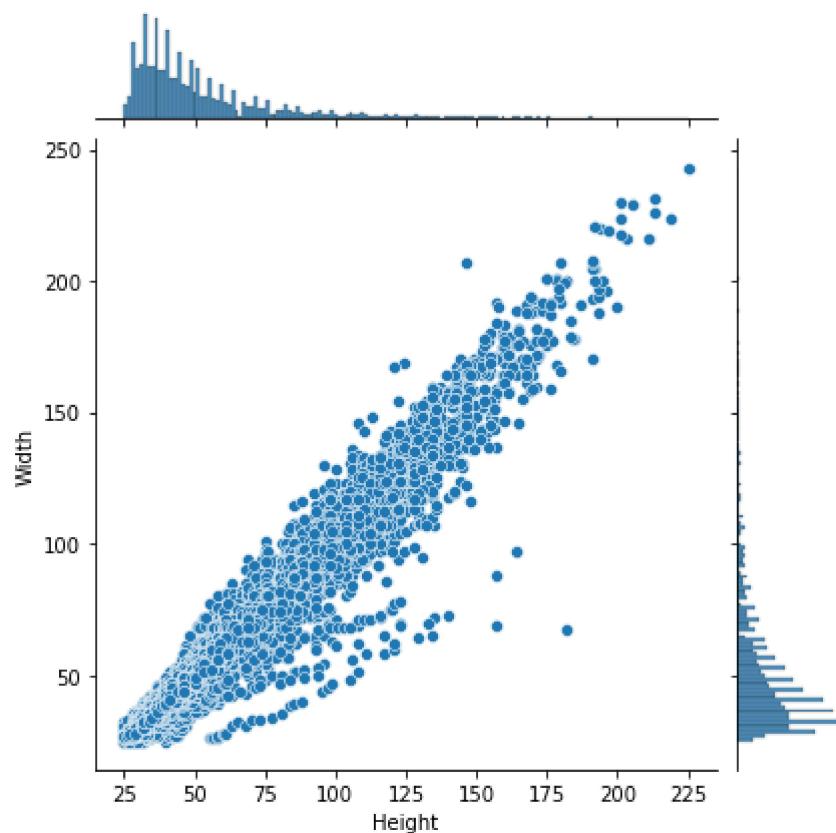
In [30]:

```
1 plt.figure(figsize=(15,8))
2 plt.xticks(rotation=70)
3 plt_box=sns.kdeplot(data=df_train, x="Height", y="Width")
```



```
In [31]: 1 sns.jointplot(data=df_train, x="Height", y="Width")
```

```
Out[31]: <seaborn.axisgrid.JointGrid at 0x1de6af4dc70>
```



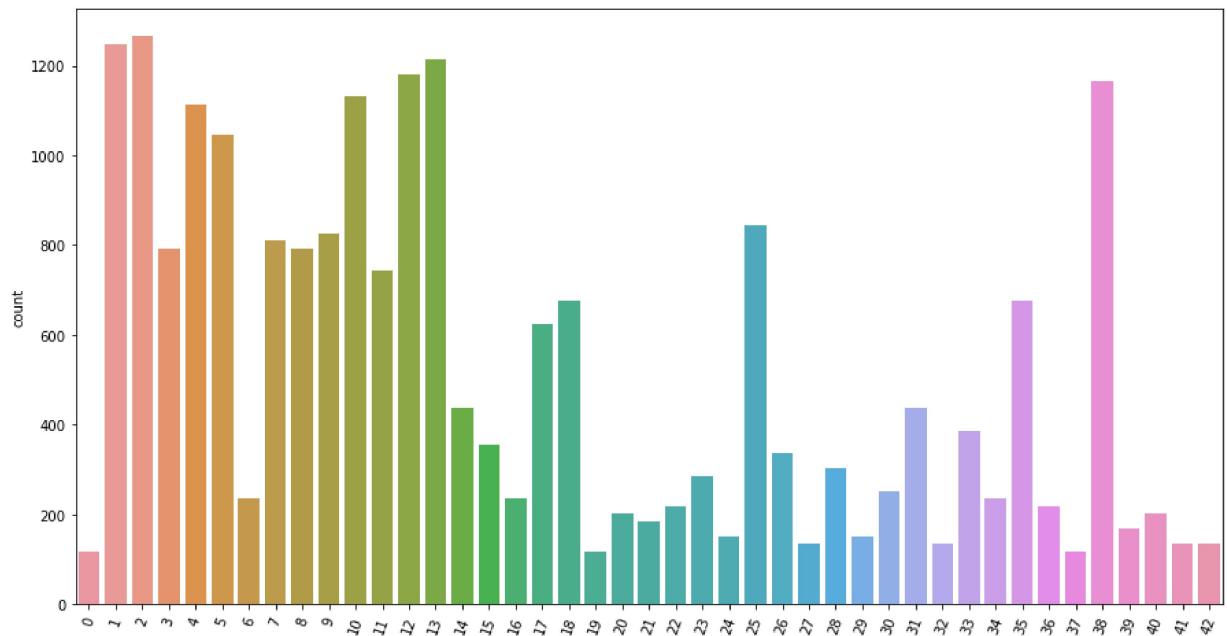
```
In [32]: 1 data = []
2 labels = []
3 classes = 43
4 #cur_path = 'C:\\\\Users\\\\adari\\\\data\\\\'
5 cur_path = 'C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Analytics\\\\Final
6 # cur_path = 'C:\\\\Users\\\\Triveni\\\\Desktop\\\\MSDSBA\\\\Winter 2021\\\\IA\\\\Project\\\\'
```

```
In [33]: 1 #Retrieving the images and their labels
2 for i in range(classes):
3     path = os.path.join(cur_path, 'Train', str(i))
4     images = os.listdir(path)
5     for a in images:
6         try:
7             image = Image.open(path + '\\\\' + a)
8             image = image.resize((30,30))
9             image = np.array(image)
10            #sim = Image.fromarray(image)
11            data.append(image)
12            labels.append(i)
13        except:
14            print("Error loading image")
```

```
In [34]: 1 #Converting Lists into numpy arrays
          2 data = np.array(data)
          3 labels = np.array(labels)
```

```
In [35]: 1 #Checking for class imbalance in training set
          2 # Below plot shows the class distribution of the training images.
          3
          4 plt.figure(figsize=(15,8))
          5 plt.xticks(rotation=70)
          6 plt_box=sns.countplot(labels)
```

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.



Observations :

1. The target classes are clearly not uniformly distributed.
2. This is quite logical since some signs like "Keep speed below 30 K-mph" or "A bump ahead" appears more often than signs like "Road under construction ahead".
3. Images are unevenly distributed among each class. Some classes have around 2500 images while others have as low as 250 images.
4. Less number of images can undermine the training process. We can use Data-Augmentation to overcome the lack of adequate number of images in some classes.

It's always a good practice to understand where your model stands without doing any preprocessing as that would help you establish a score for your model, which you could improve upon each iteration. The evaluation metric for our model would be "accuracy" score. Using a fully connected Sequential Neural Network Architecture for Baseline Scores and other testing

ESTABLISHING SCORE WITHOUT ANY

PREPROCESSING

The model = Sequential() statements loads the network. The input shape is 2700 (as images have 3 color channels). The Activation function is “relu”. During hyperparameters optimization we can check with Tanh, Sigmoid and other activation function if they are better suited for the task. For now we stick on to “relu”. There are 4 hidden layers of 128 neurons with relu activation and after each hidden layer except the last one a dropout(50%) function is included. The output layer has the softmax activation since we are dealing with multi class classification and there are 43 classes.

```
In [36]: 1 #Splitting training and testing dataset  
2 X_train,X_val, y_train, y_val = train_test_split(data, labels, test_size=0.2)
```

```
In [37]: 1 from sklearn.metrics import accuracy_score  
2 # y_test = pd.read_csv('C:\\Users\\Triveni\\Desktop\\MSDSBA\\Winter 2021\\IA  
3 y_test = pd.read_csv('C:\\Drive\\Siri\\MSDSBA\\4. IE7860 Intelligent Analyti  
4 #y_test = pd.read_csv('C:\\Users\\adari\\data\\Test.csv')  
5  
6 labels = y_test["ClassId"].values  
7 imgs = y_test["Path"].values  
8 data=[]  
9 for img in imgs:  
10    image = Image.open("C:\\Users\\Triveni\\Desktop\\MSDSBA\\Winter 2021\\  
11       image = Image.open("C:\\Drive\\Siri\\MSDSBA\\4. IE7860 Intelligent Analy  
12       #image = Image.open("C:\\Users\\adari\\data\\\"+img)  
13       image = image.resize((30,30))  
14       data.append(np.array(image))  
15 X_test=np.array(data)
```

```
In [38]: 1 #Converting the Labels into one hot encoding  
2  
3 y_train_cnn = to_categorical(y_train, 43)  
4 y_val_cnn = to_categorical(y_val, 43)
```

```
In [39]: 1 y_train_cnn.shape
```

```
Out[39]: (17625, 43)
```

```
In [40]: 1 X_train=np.array(X_train)
```

```
In [41]: 1 X_val=np.array(X_val)
```

```
In [42]: 1 print(X_train.shape)  
2 print(y_train.shape)
```

```
(17625, 30, 30, 3)  
(17625,)
```

```
In [43]: 1 print(X_val.shape)
          2 print(y_val.shape)
```

```
(4407, 30, 30, 3)
(4407,)
```

```
In [44]: 1 X_test.shape
```

```
Out[44]: (12630, 30, 30, 3)
```

```
In [45]: 1 y_test.shape
```

```
Out[45]: (12630, 8)
```

```
In [46]: 1 # # Flatten input: Changing dimension of input images from N*30*30*3 to N*27
          2 X_train_rs = X_train.reshape([X_train.shape[0], -1])
          3 X_val_rs = X_val.reshape([X_val.shape[0], -1])
```

```
In [47]: 1 print(X_train_rs.shape)
          2 print(X_val_rs.shape)
```

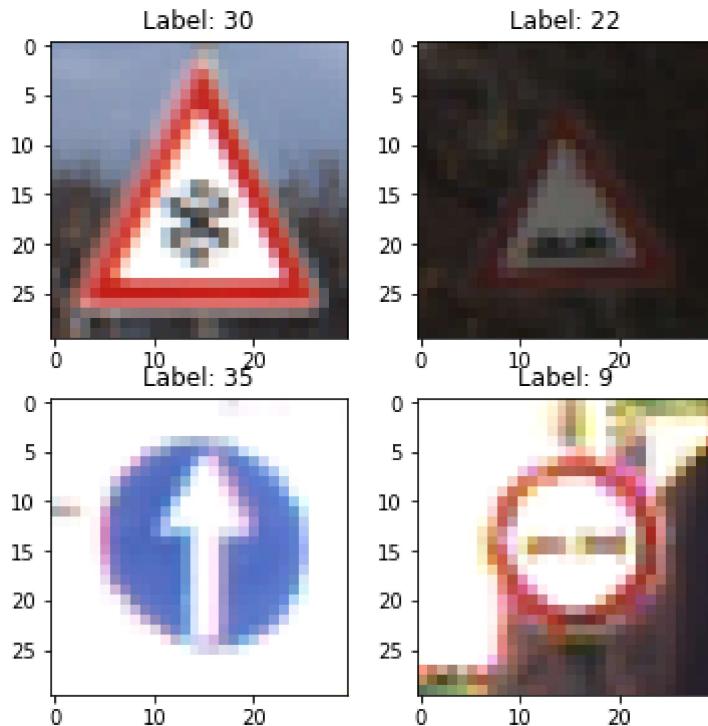
```
(17625, 2700)
(4407, 2700)
```

```
In [48]: 1 X_test_rs = X_test.reshape([X_test.shape[0], -1])
```

```
In [49]: 1 print(X_test_rs.shape)
```

```
(12630, 2700)
```

```
In [50]: 1 plt.figure(figsize=[6,6])
2 for i in range(4):
3     plt.subplot(2,2,i+1)
4     plt.title("Label: %i"%y_train[i])
5     plt.imshow(X_train_rs[i].reshape([30,30,3]));
```



```
In [51]: 1 ## Changing Labels to one-hot encoded vector
2 from sklearn.preprocessing import LabelBinarizer
3 lb = LabelBinarizer()
4 y_train_one_hot = lb.fit_transform(y_train)
5 y_val_one_hot = lb.transform(y_val)
6 print('Train labels dimension:');print(y_train_one_hot.shape)
7 print('Test labels dimension:');print(y_val_one_hot.shape)
```

Train labels dimension:
(17625, 43)
Test labels dimension:
(4407, 43)

```
In [52]: 1 from sklearn.preprocessing import LabelBinarizer
2 lb = LabelBinarizer()
3 y_test_one_hot = lb.fit_transform(labels)
4 print('Train labels dimension:');print(y_train_one_hot.shape)
5 print('Test labels dimension:');print(y_val_one_hot.shape)
```

Train labels dimension:
(17625, 43)
Test labels dimension:
(4407, 43)

Basic Models

MLP

In [53]:

```
1 from keras.models import Sequential
2 from tensorflow.keras.layers import BatchNormalization
3 n_classes = len(np.unique(labels))
4 model=tf.keras.Sequential()
5 model.add(keras.layers.Dense(128,input_dim=2700, activation='relu'))
6 model.add(BatchNormalization())
7 model.add(keras.layers.Dense(128,activation='relu'))
8 model.add(BatchNormalization())
9 model.add(keras.layers.Dropout(0.5))
10 model.add(keras.layers.Dense(128,activation='relu'))
11 model.add(BatchNormalization())
12 model.add(keras.layers.Dropout(0.5))
13 model.add(keras.layers.Dense(128,activation='relu'))
14 model.add(BatchNormalization())
15 model.add(keras.layers.Dense(n_classes,activation='softmax'))
16 model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=[ 'a
```

In [54]:

```
1 history = model.fit(X_train_rs, y_train_one_hot, validation_data = (X_val_rs,
2 score_mlp = model.evaluate(X_val_rs, y_val_one_hot,batch_size=1, verbose=0)
3 print('Test loss:', score_mlp[0], ' Test accuracy:', score_mlp[1])
4
```

```
Epoch 1/15
138/138 - 3s - loss: 3.2831 - accuracy: 0.1750 - val_loss: 2.7257 - val_accuracy: 0.2539
Epoch 2/15
138/138 - 1s - loss: 2.0110 - accuracy: 0.4137 - val_loss: 1.9887 - val_accuracy: 0.4118
Epoch 3/15
138/138 - 1s - loss: 1.4596 - accuracy: 0.5540 - val_loss: 1.3704 - val_accuracy: 0.6031
Epoch 4/15
138/138 - 1s - loss: 1.1062 - accuracy: 0.6637 - val_loss: 1.0502 - val_accuracy: 0.6742
Epoch 5/15
138/138 - 1s - loss: 0.9230 - accuracy: 0.7183 - val_loss: 0.8034 - val_accuracy: 0.7497
Epoch 6/15
138/138 - 1s - loss: 0.8134 - accuracy: 0.7525 - val_loss: 0.6515 - val_accuracy: 0.8003
Epoch 7/15
138/138 - 1s - loss: 0.7216 - accuracy: 0.7809 - val_loss: 0.6030 - val_accuracy: 0.8273
Epoch 8/15
138/138 - 1s - loss: 0.6890 - accuracy: 0.7918 - val_loss: 0.6609 - val_accuracy: 0.8046
Epoch 9/15
138/138 - 1s - loss: 0.5995 - accuracy: 0.8179 - val_loss: 0.5922 - val_accuracy: 0.8353
Epoch 10/15
138/138 - 1s - loss: 0.5772 - accuracy: 0.8237 - val_loss: 0.6839 - val_accuracy: 0.7910
Epoch 11/15
138/138 - 1s - loss: 0.5403 - accuracy: 0.8376 - val_loss: 0.5348 - val_accuracy: 0.8448
Epoch 12/15
138/138 - 1s - loss: 0.5106 - accuracy: 0.8441 - val_loss: 0.9075 - val_accuracy: 0.7663
Epoch 13/15
138/138 - 1s - loss: 0.4969 - accuracy: 0.8505 - val_loss: 0.5267 - val_accuracy: 0.8475
Epoch 14/15
138/138 - 1s - loss: 0.4569 - accuracy: 0.8592 - val_loss: 0.3782 - val_accuracy: 0.8902
Epoch 15/15
138/138 - 1s - loss: 0.4340 - accuracy: 0.8673 - val_loss: 0.5083 - val_accuracy: 0.8500
Test loss: 0.5083445906639099      Test accuracy: 0.8500113487243652
```

```
In [55]: 1 predict_mlp = model.predict(X_test_rs, batch_size = 1)
2 score_test_mlp=model.evaluate(X_test_rs, y_test_one_hot,batch_size=1, verbose=0)
3
```

```
In [56]: 1 score_test_mlp[1]
```

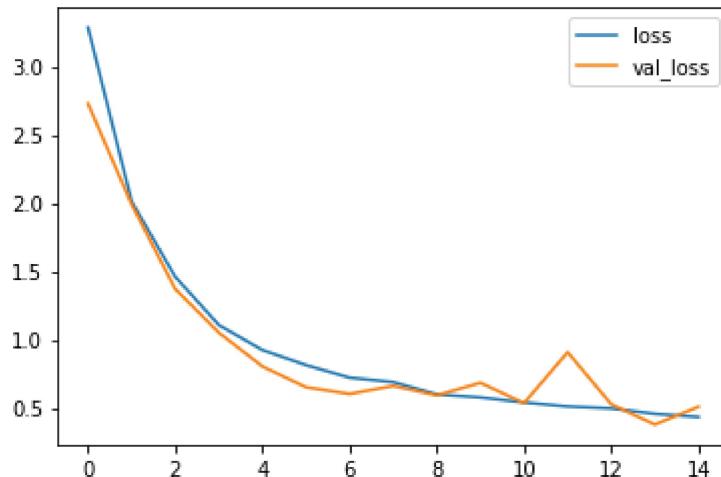
```
Out[56]: 0.7820269465446472
```

```
In [57]: 1
2 for name, value in zip(model.metrics_names, score_test_mlp):
3     print(name, ': ', value)
4
5 print()
6
7
```

```
loss :  0.8219485878944397
accuracy :  0.7820269465446472
```

```
In [58]: 1 # Leraning curve MLP
2 plt.plot(history.history['loss'])
3 plt.plot(history.history['val_loss'])
4 plt.legend(['loss','val_loss'])
```

```
Out[58]: <matplotlib.legend.Legend at 0x1de77a24b50>
```



```
In [59]: 1 val_performance = {}
2 performance = {}
3 Model={} 
```

```
In [60]: 1 val_performance['MLP']=score_mlp[1]
2 performance['MLP']=score_test_mlp[1]
3 Model['MLP']='MLP'
```

The model was able to achieve an accuracy score of 83.61% without any preprocessing.

Base cnn model

In [61]:

```
1  cnn_model = tf.keras.Sequential([
2      keras.layers.Conv2D(filters=32,kernel_size=(5,5),activation='relu',input_
3          keras.layers.MaxPool2D(pool_size=(2,2)),# down sampling the output inste_
4              keras.layers.Dropout(0.2),
5                  keras.layers.Flatten(), # flatten out the layers
6                      keras.layers.Dense(256,activation='relu'),
7                          keras.layers.Dense(43,activation = 'softmax')
8
9  ])
10 #Compilation of the model
11 cnn_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics_
12 epochs = 15
13
```

```
In [62]: 1 history = cnn_model.fit(X_train, y_train_cnn, epochs=epochs, validation_dat
Epoch 1/15
551/551 [=====] - 14s 24ms/step - loss: 9.8534 - accuracy: 0.4055 - val_loss: 0.6527 - val_accuracy: 0.8328
Epoch 2/15
551/551 [=====] - 13s 24ms/step - loss: 0.6735 - accuracy: 0.8282 - val_loss: 0.6572 - val_accuracy: 0.8514
Epoch 3/15
551/551 [=====] - 13s 24ms/step - loss: 0.4914 - accuracy: 0.8728 - val_loss: 0.5357 - val_accuracy: 0.8716
Epoch 4/15
551/551 [=====] - 14s 25ms/step - loss: 0.5386 - accuracy: 0.8671 - val_loss: 0.3730 - val_accuracy: 0.9115
Epoch 5/15
551/551 [=====] - 14s 26ms/step - loss: 0.3917 - accuracy: 0.8996 - val_loss: 0.4602 - val_accuracy: 0.8981
Epoch 6/15
551/551 [=====] - 13s 24ms/step - loss: 0.3389 - accuracy: 0.9139 - val_loss: 0.3993 - val_accuracy: 0.9247
Epoch 7/15
551/551 [=====] - 13s 24ms/step - loss: 0.4064 - accuracy: 0.9073 - val_loss: 0.3226 - val_accuracy: 0.9369
Epoch 8/15
551/551 [=====] - 14s 25ms/step - loss: 0.3210 - accuracy: 0.9261 - val_loss: 0.3985 - val_accuracy: 0.9238
Epoch 9/15
551/551 [=====] - 14s 25ms/step - loss: 0.3574 - accuracy: 0.9224 - val_loss: 0.4729 - val_accuracy: 0.9158
Epoch 10/15
551/551 [=====] - 14s 26ms/step - loss: 0.2975 - accuracy: 0.9346 - val_loss: 0.3267 - val_accuracy: 0.9412
Epoch 11/15
551/551 [=====] - 13s 24ms/step - loss: 0.2475 - accuracy: 0.9397 - val_loss: 0.7539 - val_accuracy: 0.8834
Epoch 12/15
551/551 [=====] - 13s 24ms/step - loss: 0.3863 - accuracy: 0.9184 - val_loss: 0.3818 - val_accuracy: 0.9401
Epoch 13/15
551/551 [=====] - 14s 25ms/step - loss: 0.2913 - accuracy: 0.9361 - val_loss: 0.4450 - val_accuracy: 0.9365
Epoch 14/15
551/551 [=====] - 14s 25ms/step - loss: 0.2408 - accuracy: 0.9457 - val_loss: 0.4388 - val_accuracy: 0.9310
Epoch 15/15
551/551 [=====] - 13s 24ms/step - loss: 0.2692 - accuracy: 0.9425 - val_loss: 0.4241 - val_accuracy: 0.9437
```

```
In [63]: 1 score_cnn = cnn_model.evaluate(X_val, y_val_cnn,batch_size=1, verbose=0)
2 print('Test loss:', score_cnn[0], ' Test accuracy:', score_cnn[1])
```

Test loss: 0.42412441968917847 Test accuracy: 0.9437258839607239

In [64]:

```
1 pred_cnn = cnn_model.predict_classes(X_test)
2 #Accuracy with the test data
3 from sklearn.metrics import accuracy_score
4 print(accuracy_score(labels, pred_cnn))
```

`model.predict_classes()` is deprecated and will be removed after 2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)` , if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")` , if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).

0.8524940617577197

In [65]:

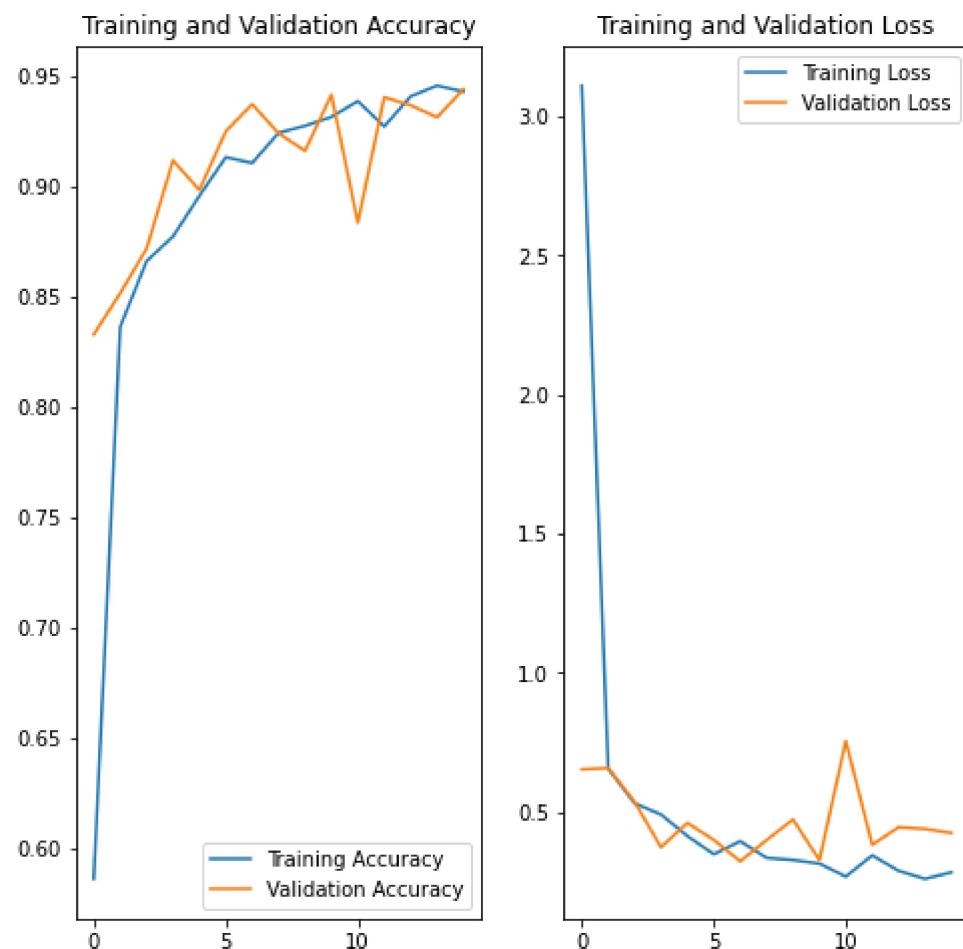
```
1 val_performance['CNN']= score_cnn[1]
2 performance['CNN']=accuracy_score(labels, pred_cnn)
3 Model['CNN']='CNN'
```

In [66]:

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6
7
```

In [67]:

```
1 epochs_range = range(15)
2
3 plt.figure(figsize=(8, 8))
4 plt.subplot(1, 2, 1)
5 plt.plot(epochs_range, acc, label='Training Accuracy')
6 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
7 plt.legend(loc='lower right')
8 plt.title('Training and Validation Accuracy')
9
10 plt.subplot(1, 2, 2)
11 plt.plot(epochs_range, loss, label='Training Loss')
12 plt.plot(epochs_range, val_loss, label='Validation Loss')
13 plt.legend(loc='upper right')
14 plt.title('Training and Validation Loss')
15 plt.show()
```



SVM

In [68]:

```
1 #Normalize the data for SVM
2 from sklearn.preprocessing import StandardScaler
3 sc = StandardScaler()
4 X_train_svm = sc.fit_transform(X_train_rs)
5 X_val_svm = sc.fit_transform(X_val_rs)
6 X_test_svm = sc.fit_transform(X_test_rs)
```

```
In [69]: 1 from sklearn.svm import SVC  
2 svclassifier = SVC(kernel='linear',decision_function_shape='ovo')  
3 #svclassifier = SVC(gamma=1,kernel='rbf')  
4 svclassifier.fit(X_train_rs, y_train)  
5 score_train = svclassifier.score(X_val_rs, y_val)  
6 Pred_svm=svclassifier.predict(X_test_rs)      #prediction
```

```
In [70]: 1 score_train
```

```
Out[70]: 0.9355570683004312
```

In [71]:

```
1 # Evaluate predictions
2 print(accuracy_score(labels, Pred_svm))
3 print(classification_report(labels, Pred_svm))
```

0.7991290577988915

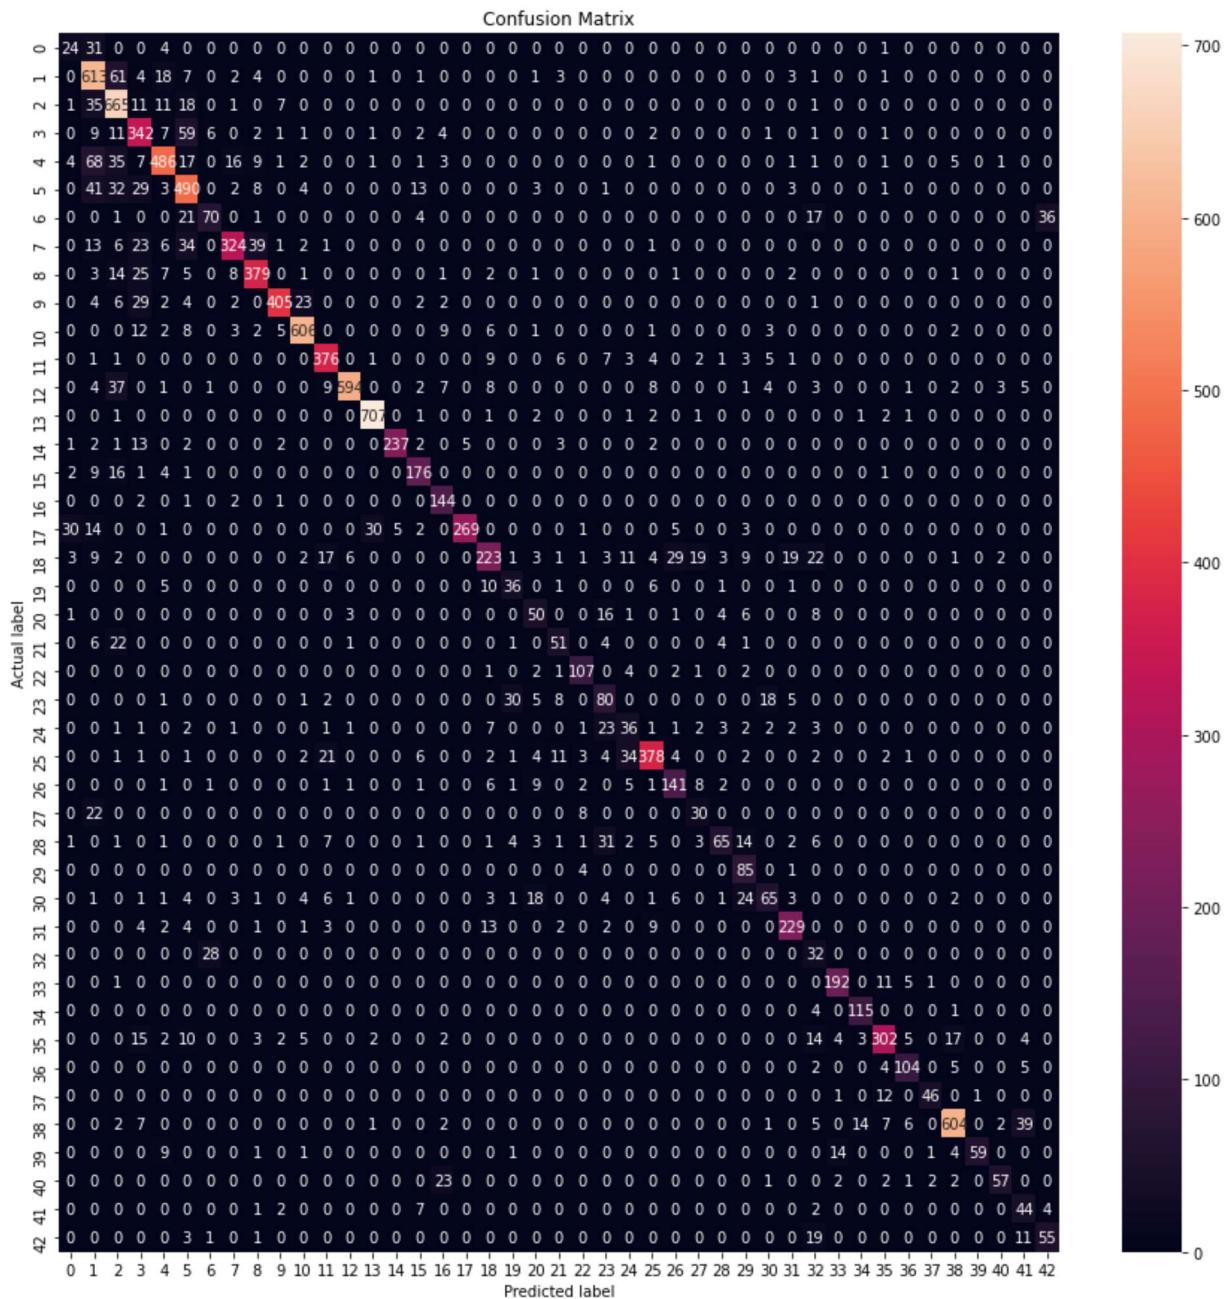
	precision	recall	f1-score	support
0	0.36	0.40	0.38	60
1	0.69	0.85	0.76	720
2	0.73	0.89	0.80	750
3	0.65	0.76	0.70	450
4	0.85	0.74	0.79	660
5	0.71	0.78	0.74	630
6	0.65	0.47	0.54	150
7	0.89	0.72	0.80	450
8	0.84	0.84	0.84	450
9	0.95	0.84	0.89	480
10	0.93	0.92	0.92	660
11	0.85	0.90	0.87	420
12	0.98	0.86	0.92	690
13	0.95	0.98	0.97	720
14	0.98	0.88	0.93	270
15	0.80	0.84	0.82	210
16	0.73	0.96	0.83	150
17	0.98	0.75	0.85	360
18	0.76	0.57	0.65	390
19	0.47	0.60	0.53	60
20	0.49	0.56	0.52	90
21	0.58	0.57	0.57	90
22	0.84	0.89	0.86	120
23	0.46	0.53	0.49	150
24	0.37	0.40	0.39	90
25	0.89	0.79	0.83	480
26	0.74	0.78	0.76	180
27	0.45	0.50	0.48	60
28	0.77	0.43	0.56	150
29	0.56	0.94	0.70	90
30	0.65	0.43	0.52	150
31	0.84	0.85	0.85	270
32	0.22	0.53	0.31	60
33	0.90	0.91	0.91	210
34	0.86	0.96	0.91	120
35	0.87	0.77	0.82	390
36	0.84	0.87	0.85	120
37	0.92	0.77	0.84	60
38	0.93	0.88	0.90	690
39	0.98	0.66	0.79	90
40	0.88	0.63	0.74	90
41	0.41	0.73	0.52	60
42	0.58	0.61	0.59	90
accuracy			0.80	12630
macro avg	0.74	0.73	0.73	12630
weighted avg	0.82	0.80	0.80	12630

In [72]:

```
1 #plot for confusion matrix
2 def plot_cm(labels, predictions):
3     cm = confusion_matrix(labels, predictions)
4     plt.figure(figsize=(15,15))
5     sns.heatmap(cm, annot=True, fmt="d")
6     plt.title('Confusion Matrix')
7     plt.ylabel('Actual label')
8     plt.xlabel('Predicted label')
```

In [73]:

```
1 plot_cm(labels, Pred_svm)
```



In [74]:

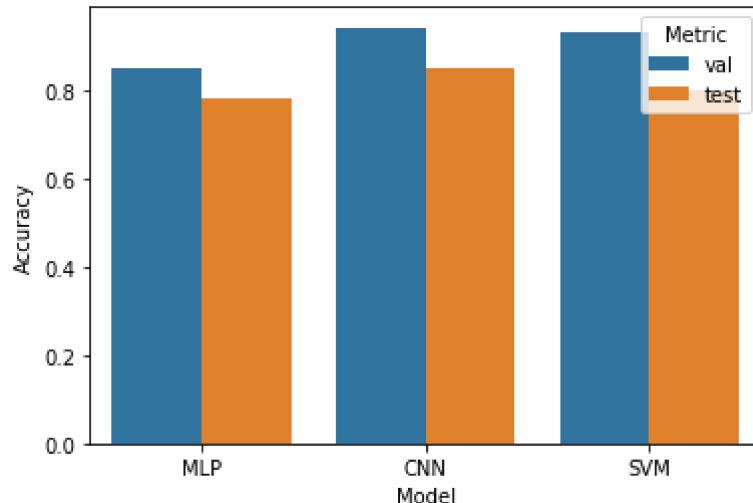
```
1 val_performance['SVM']= score_train
2 performance['SVM']=accuracy_score(labels, Pred_svm)
3 Model['SVM']='SVM'
```

Model Performance

```
In [75]: 1 Performance = pd.DataFrame(  
2     {'val': val_performance,  
3      'test': performance,  
4      'Model':Model  
5 })
```

```
In [76]: 1 df_accuracy = Performance.melt(id_vars=['Model'], var_name='Metric', value_n
```

```
In [77]: 1 ax = sns.barplot(x="Model", y="Accuracy", hue="Metric", data=df_accuracy)
```



PRE-PROCESSING STEPS

```
In [78]: 1 print("Number of training examples =", df_train.shape[0])  
2 print("Number of testing examples =", df_test.shape[0])  
3  
4 print("Number of classes =", y_val_one_hot.shape[1])
```

```
Number of training examples = 39209  
Number of testing examples = 12630  
Number of classes = 43
```

Preprocess the data here. Preprocessing steps could include normalization, converting to grayscale, etc.

Equalize histograms of training samples - by generation of additional, transformed images

In [79]:

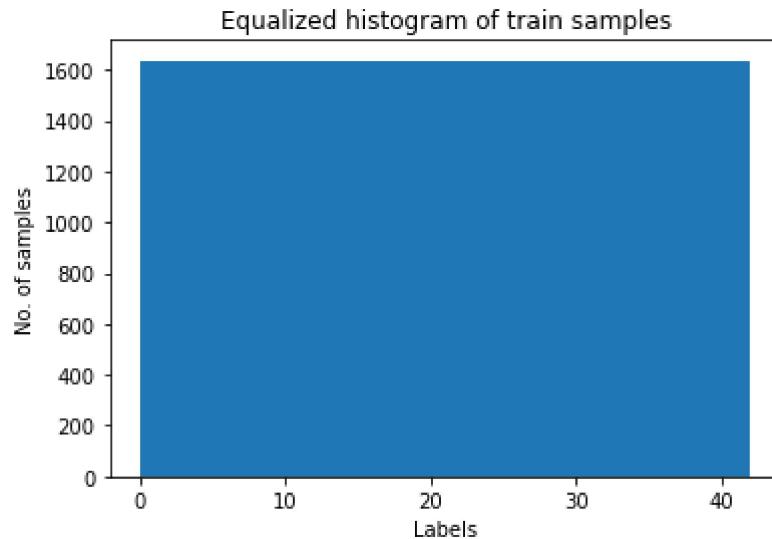
```
1 from tqdm import tqdm
2 from scipy import ndimage
3 import cv2
4
5 def augment_brightness_camera_images(image):
6     image1 = cv2.cvtColor(image, cv2.COLOR_RGB2HSV)
7     random_bright = .25+np.random.uniform()
8     image1[:, :, 2] = image1[:, :, 2]*random_bright
9     image1 = cv2.cvtColor(image1, cv2.COLOR_HSV2RGB)
10    return image1
11
12 def transform_image(img):
13     ang_range = 25
14     ang_rot = np.random.uniform(ang_range)-ang_range/2
15     rows, cols, ch = img.shape
16     Rot_M = cv2.getRotationMatrix2D((cols/2, rows/2), ang_rot, 1)
17
18     img = cv2.warpAffine(img, Rot_M, (cols, rows))
19     img = augment_brightness_camera_images(img)
20
21    return img
22
23 def get_random_image_of_given_label(images_set, labels_set, label):
24     image_indexes = np.where(labels_set == label)
25     rand_index = random.randint(0, np.bincount(labels_set)[label] - 1)
26     return images_set[image_indexes][rand_index]
27
28 def equalize_samples_set(X_set, y_set):
29     labels_count_arr = np.bincount(y_set)
30     labels_bins = np.arange(len(labels_count_arr))
31
32     ind = 0
33
34     for label in tqdm(labels_bins):
35         labels_no_to_add = int(np.mean(labels_count_arr)) * 4 - labels_coun
36
37         ind = ind + 1
38         X_temp = []
39         y_temp = []
40
41         for num in range(labels_no_to_add):
42             rand_image = get_random_image_of_given_label(X_set, y_set, label)
43             X_temp.append(transform_image(rand_image))
44             y_temp.append(label)
45
46         X_set = np.append(X_set, np.array(X_temp), axis=0)
47         y_set = np.append(y_set, np.array(y_temp), axis=0)
48
49    return X_set, y_set
50
```

In [80]:

```
1 n_train = X_train.shape[0]
2 X_train_eq, y_train_eq = equalize_samples_set(X_train, y_train)
3
4
5 n_classes = len(np.unique(labels))
6 n, bins, patches = plt.hist(y_train_eq, n_classes)
7 plt.xlabel('Labels')
8 plt.ylabel('No. of samples')
9 plt.title('Equalized histogram of train samples')
10 plt.show()
11
12 print("Train set increased from {} to {}".format(n_train,X_train_eq.shape[0])
13
14 fig=plt.figure()
15 fig.suptitle('Test image to transform', fontsize=16)
16 plt.imshow(X_train_eq[1000])
17
18 grid_len = 4
19 fig=plt.figure(figsize=(grid_len,grid_len))
20 fig.suptitle('Test images after transformation', fontsize=16)
21
22 for i in range(1,grid_len*grid_len+1):
23     image = transform_image(X_train_eq[1000])
24     plt.subplot(grid_len,grid_len,i)
25     plt.imshow(image)
```

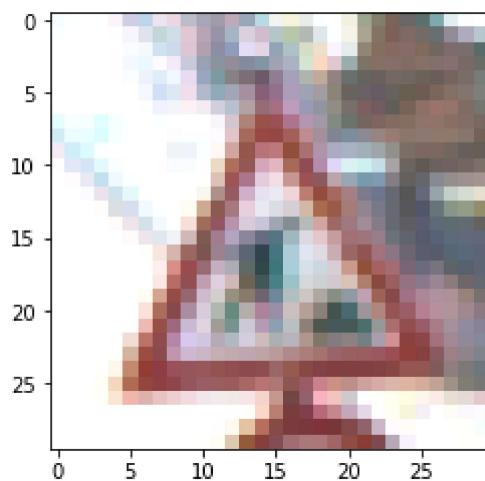
100% |

| 43/43 [00:26<00:00, 1.64it/s]



Train set increased from 17625 to 70348

Test image to transform



Test images after transformation

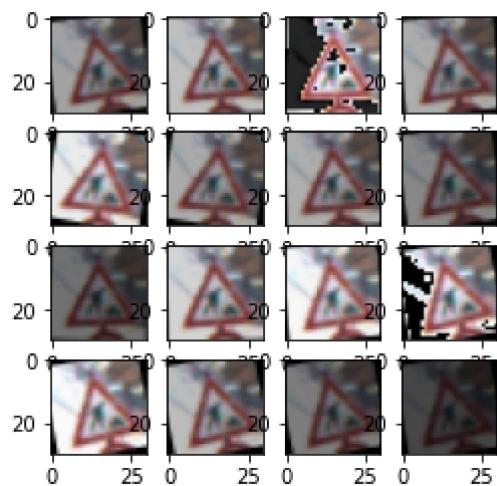


Image normalization and grayscale conversion

In [81]:

```
1 import cv2
2 import tensorflow as tf
3 # from tensorflow.contrib.layers import flatten
4
5 def grayscale(img):
6     return cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)[:, :, None]
7
8 def normalize(value):
9     return value / 255
10
11 def preprocess_image(image):
12     #img = grayscale(image)
13     img = normalize(image)
14     return img
15
16 from tqdm import tqdm
17 def preprocess_batch(images):
18     imgs = np.zeros(shape=images.shape)
19     processed_image_depth = preprocess_image(images[0]).shape[2]
20     imgs = imgs[:, :, :, 0:processed_image_depth]
21     for i in tqdm(range(images.shape[0])):
22         imgs[i] = preprocess_image(images[i])
23     return imgs
24
25
26 # X_train_eq, y_train_eq
27
28 X_train_processed = preprocess_batch(X_train_eq)
29 X_val_processed = preprocess_batch(X_val)
30
31 no_test_image = 1000
32
33 sample_image = X_train[no_test_image]
34 sample_image_processed = normalize(X_train_eq[no_test_image])
35 fig=plt.figure(figsize=(16,3))
36 sub=plt.subplot(131)
37 sub.set_title("Original image")
38 plt.imshow(sample_image)
39 sub=plt.subplot(132)
40 sub.set_title("Preprocessed image")
41 plt.imshow(sample_image_processed.squeeze(), cmap='gray')
42
43 print("Sample image dimension BEFORE processing: {}".format(sample_image.sha
44 print("Sample image dimension AFTER processing: {}".format(sample_image_proc
45
46 image_depth = X_train_processed.shape[3]
47
48 sample_image = X_train[no_test_image]
49 dim1 = sample_image.shape[0]
50 dim2 = sample_image.shape[1]
51 dim3 = sample_image.shape[2]
52 sample_image_reshaped = np.reshape(sample_image, dim1*dim2*dim3)
53 plt.figure(figsize=(16,3))
54 sub=plt.subplot(131)
55 sub.set_title("Original image histogram")
56 n, bins, patches = plt.hist(sample_image_reshaped, 255)
```

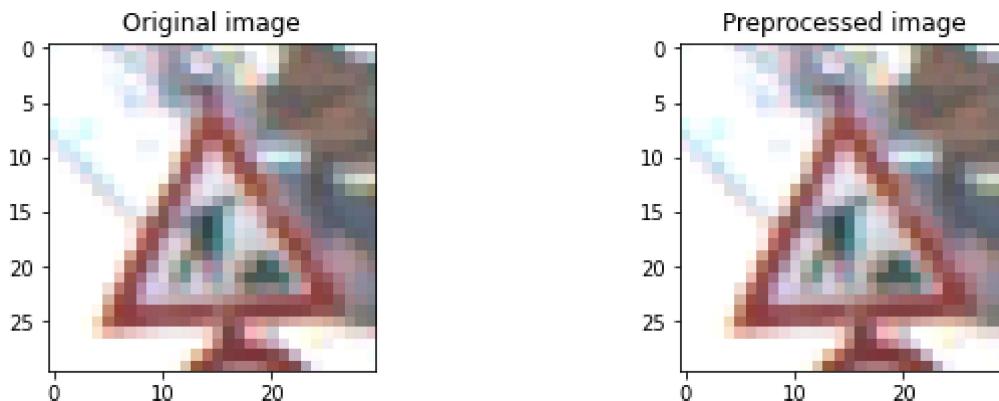
```

57
58 sample_image_processed = X_train_processed[no_test_image]
59 dim1 = sample_image_processed.shape[0]
60 dim2 = sample_image_processed.shape[1]
61 dim3 = sample_image_processed.shape[2]
62 sample_image_processed_reshaped = np.reshape(sample_image_processed, dim1*di
63 sub=plt.subplot(132)
64 sub.set_title("Preprocessed image histogram")
65 n, bins, patches = plt.hist(sample_image_processed_reshaped, 255)
100%|██████████| 70348/70348 [00:00<00:00, 75661.98it/s]
100%|██████████| 4407/4407 [00:00<00:00, 73640.27it/s]

```

Sample image dimension BEFORE processing: (30, 30, 3)

Sample image dimension AFTER processing: (30, 30, 3)



```

In [82]: 1 # # Flatten input: Changing dimension of input images from N*30*30*3 to N*27
2 X_train_pr = X_train_processed.reshape([X_train_processed.shape[0], -1])
3 X_val_pr = X_val_processed.reshape([X_val_processed.shape[0], -1])
4
5 lb = LabelBinarizer()
6 y_train_eq_one_hot = lb.fit_transform(y_train_eq)
7

```

```
In [83]: 1 X_train_pr.shape
```

Out[83]: (70348, 2700)

```
In [84]: 1 y_train_eq_one_hot.shape
```

Out[84]: (70348, 43)

```
In [85]: 1 y_val_one_hot.shape
```

Out[85]: (4407, 43)

```
In [86]: 1 X_val_pr.shape
```

Out[86]: (4407, 2700)

```
In [87]: 1 X_train_processed.shape
```

```
Out[87]: (70348, 30, 30, 3)
```

```
In [88]: 1 y_train_eq.shape
```

```
Out[88]: (70348,)
```

Models Improvement

SVM

```
In [89]: 1 from sklearn.svm import SVC  
2 svclassifier_pr = SVC(kernel='linear',decision_function_shape='ovo')  
3 #svclassifier = SVC(gamma=1,kernel='rbf')  
4 svclassifier_pr.fit(X_train_pr, y_train_eq)  
5  
6
```

```
Out[89]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
             decision_function_shape='ovo', degree=3, gamma='scale', kernel='linear',  
             max_iter=-1, probability=False, random_state=None, shrinking=True,  
             tol=0.001, verbose=False)
```

```
In [90]: 1 score_train_pr = svclassifier_pr.score(X_val_pr, y_val)
```

```
In [91]: 1 Pred_svm_pr=svclassifier_pr.predict(X_test_rs)      #prediction
```

```
In [92]: 1 score_train_pr
```

```
Out[92]: 0.9357839800317677
```

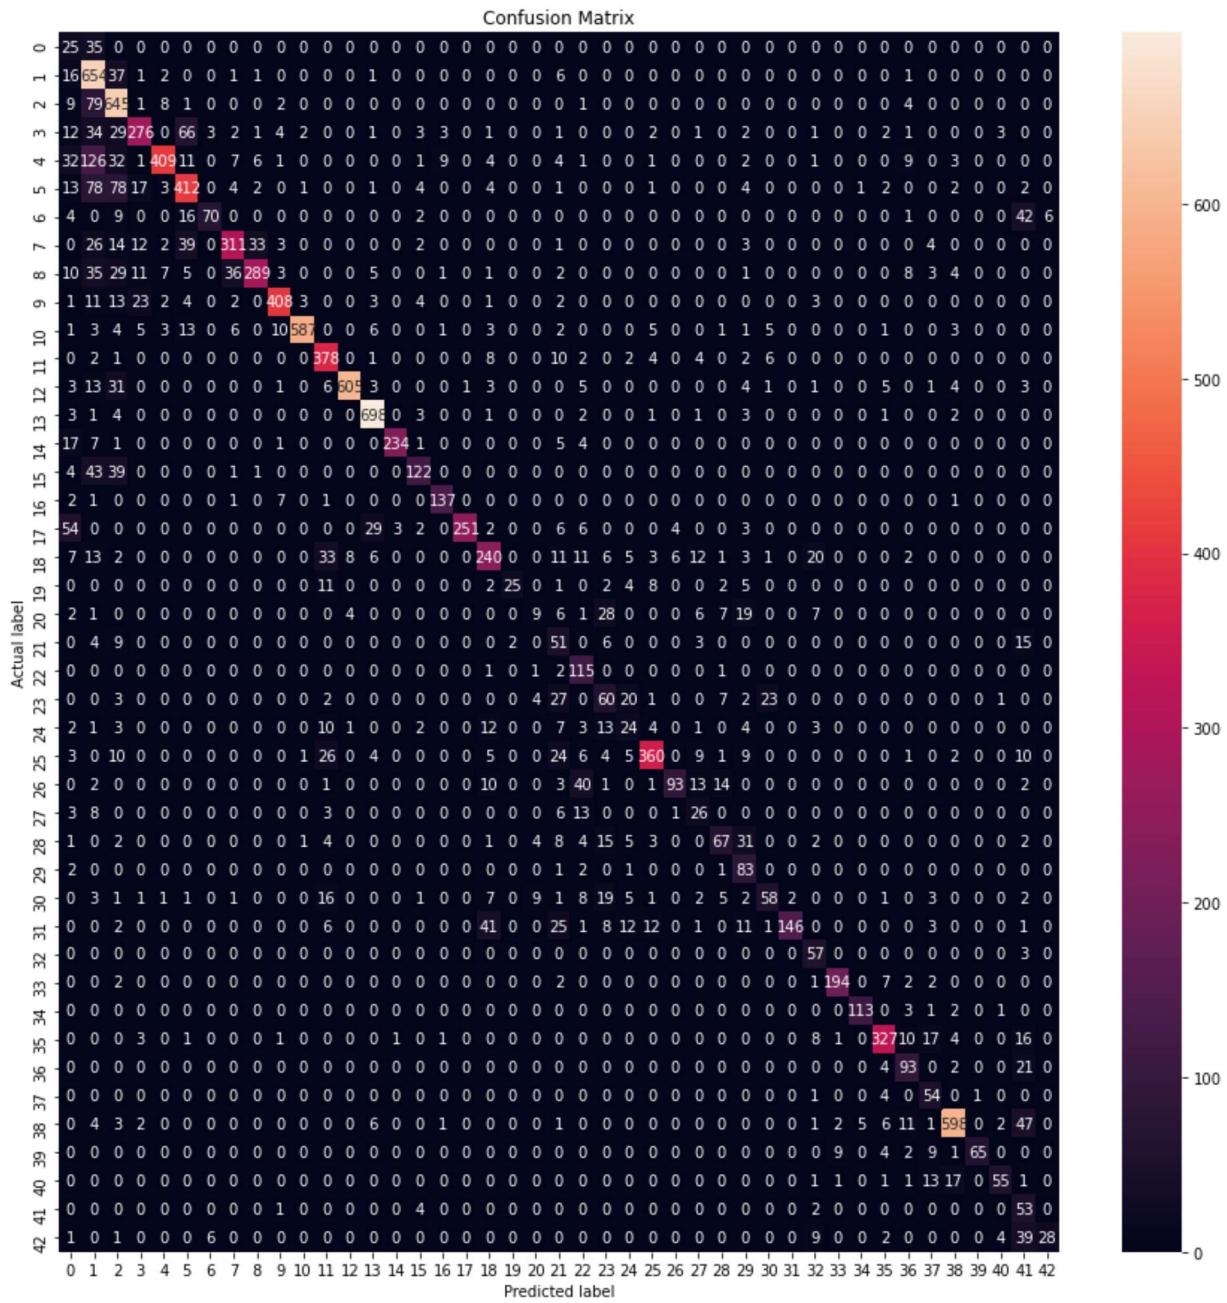
In [93]:

```
1 # Evaluate predictions
2 print(accuracy_score(labels, Pred_svm_pr))
3 print(classification_report(labels, Pred_svm_pr))
```

0.7525732383214568

	precision	recall	f1-score	support
0	0.11	0.42	0.17	60
1	0.55	0.91	0.69	720
2	0.64	0.86	0.74	750
3	0.78	0.61	0.69	450
4	0.94	0.62	0.75	660
5	0.72	0.65	0.69	630
6	0.89	0.47	0.61	150
7	0.84	0.69	0.76	450
8	0.87	0.64	0.74	450
9	0.92	0.85	0.89	480
10	0.99	0.89	0.94	660
11	0.76	0.90	0.82	420
12	0.98	0.88	0.93	690
13	0.91	0.97	0.94	720
14	0.98	0.87	0.92	270
15	0.81	0.58	0.68	210
16	0.90	0.91	0.90	150
17	1.00	0.70	0.82	360
18	0.69	0.62	0.65	390
19	0.93	0.42	0.57	60
20	0.33	0.10	0.15	90
21	0.24	0.57	0.33	90
22	0.51	0.96	0.67	120
23	0.37	0.40	0.38	150
24	0.29	0.27	0.28	90
25	0.88	0.75	0.81	480
26	0.89	0.52	0.65	180
27	0.33	0.43	0.37	60
28	0.63	0.45	0.52	150
29	0.43	0.92	0.58	90
30	0.61	0.39	0.47	150
31	0.99	0.54	0.70	270
32	0.48	0.95	0.64	60
33	0.94	0.92	0.93	210
34	0.95	0.94	0.95	120
35	0.89	0.84	0.86	390
36	0.62	0.78	0.69	120
37	0.49	0.90	0.63	60
38	0.93	0.87	0.90	690
39	0.98	0.72	0.83	90
40	0.83	0.61	0.71	90
41	0.20	0.88	0.33	60
42	0.82	0.31	0.45	90
accuracy			0.75	12630
macro avg	0.72	0.69	0.67	12630
weighted avg	0.81	0.75	0.76	12630

```
In [94]: 1 plot_cm(labels, Pred_svm_pr)
```



```
In [95]: 1 val_performance_pr = {}  
2 performance_pr = {}  
3 Model_pr={}
```

```
In [96]: 1 val_performance_pr['SVM']= score_train_pr  
2 performance_pr['SVM']=accuracy_score(labels, Pred_svm_pr)  
3 Model_pr['SVM']='SVM'
```

CNN_1

```
In [97]: 1 img_height = 60  
2 img_width = 60
```

```
In [98]: 1 y_train_prcnn = to_categorical(y_train_eq, 43)  
2 y_val_prcnn = to_categorical(y_val, 43)
```

```
In [99]: 1 X_train_processed.shape
```

```
Out[99]: (70348, 30, 30, 3)
```

```
In [100]: 1 X_train.shape[1:]
```

```
Out[100]: (30, 30, 3)
```

```
In [101]: 1 cnn_ag_model = Sequential([  
2     Conv2D(filters=16,kernel_size=(3,3),activation='relu',input_shape=(30,30  
3     MaxPool2D(pool_size=(2,2)),# down sampling the output instead of 28*28 i  
4     Conv2D(32, kernel_size=(4, 4), activation='relu'),  
5     Conv2D(32, kernel_size=(4, 4), activation='relu'),  
6     MaxPool2D(pool_size=(2, 2)),  
7     #     Conv2D(64, kernel_size=(3,3), activation='relu'),  
8     #     Conv2D(64, kernel_size=(3,3), activation='relu'),  
9     #     MaxPool2D(pool_size=(2, 2)),  
10    Dropout(0.3),  
11    Flatten(), # flatten out the layers  
12    Dense(256,activation='relu'),  
13    Dense(43,activation = 'softmax')  
14  
15    ])  
16  
17  
18
```

```
In [102]: 1 cnn_ag_model.compile(loss='categorical_crossentropy', optimizer='adam', metr
```

```
In [103]: 1 history_ag = cnn_ag_model.fit(X_train_processed, y_train_prcnn, epochs= 15,
    <   >
Epoch 1/15
2199/2199 [=====] - 43s 19ms/step - loss: 1.5164 - accuracy: 0.5836 - val_loss: 0.1213 - val_accuracy: 0.9637
Epoch 2/15
2199/2199 [=====] - 44s 20ms/step - loss: 0.2086 - accuracy: 0.9377 - val_loss: 0.0784 - val_accuracy: 0.9789
Epoch 3/15
2199/2199 [=====] - 44s 20ms/step - loss: 0.1323 - accuracy: 0.9600 - val_loss: 0.0520 - val_accuracy: 0.9850
Epoch 4/15
2199/2199 [=====] - 44s 20ms/step - loss: 0.0989 - accuracy: 0.9693 - val_loss: 0.0449 - val_accuracy: 0.9873
Epoch 5/15
2199/2199 [=====] - 45s 20ms/step - loss: 0.0824 - accuracy: 0.9751 - val_loss: 0.0429 - val_accuracy: 0.9873
Epoch 6/15
2199/2199 [=====] - 45s 21ms/step - loss: 0.0740 - accuracy: 0.9763 - val_loss: 0.0276 - val_accuracy: 0.9925
Epoch 7/15
2199/2199 [=====] - 45s 20ms/step - loss: 0.0580 - accuracy: 0.9827 - val_loss: 0.0234 - val_accuracy: 0.9943
Epoch 8/15
2199/2199 [=====] - 44s 20ms/step - loss: 0.0506 - accuracy: 0.9847 - val_loss: 0.0270 - val_accuracy: 0.9932
Epoch 9/15
2199/2199 [=====] - 45s 20ms/step - loss: 0.0452 - accuracy: 0.9857 - val_loss: 0.0266 - val_accuracy: 0.9932
Epoch 10/15
2199/2199 [=====] - 45s 20ms/step - loss: 0.0416 - accuracy: 0.9868 - val_loss: 0.0330 - val_accuracy: 0.9905
Epoch 11/15
2199/2199 [=====] - 45s 20ms/step - loss: 0.0432 - accuracy: 0.9872 - val_loss: 0.0243 - val_accuracy: 0.9946
Epoch 12/15
2199/2199 [=====] - 44s 20ms/step - loss: 0.0426 - accuracy: 0.9874 - val_loss: 0.0234 - val_accuracy: 0.9941
Epoch 13/15
2199/2199 [=====] - 45s 20ms/step - loss: 0.0400 - accuracy: 0.9870 - val_loss: 0.0312 - val_accuracy: 0.9941
Epoch 14/15
2199/2199 [=====] - 46s 21ms/step - loss: 0.0406 - accuracy: 0.9879 - val_loss: 0.0293 - val_accuracy: 0.9925
Epoch 15/15
2199/2199 [=====] - 45s 21ms/step - loss: 0.0371 - accuracy: 0.9893 - val_loss: 0.0243 - val_accuracy: 0.9948
```

```
In [104]: 1 score_cnn_pr2 = cnn_ag_model.evaluate(X_val, y_val_cnn,batch_size=1, verbose=0)
2 print('Test loss:', score_cnn_pr2[0], ' Test accuracy:', score_cnn_pr2[1])
```

Test loss: 8.642882347106934 Test accuracy: 0.9779895544052124

In [105]:

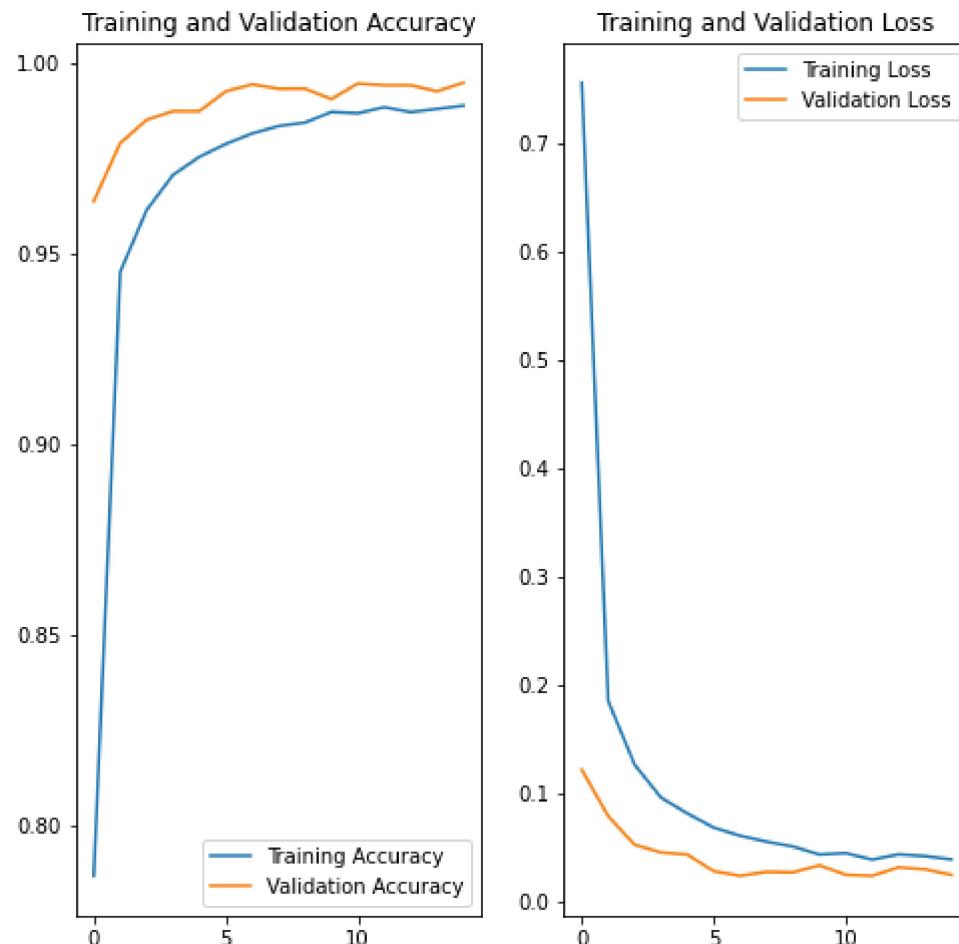
```
1 pred_cnn_pr2 = cnn_ag_model.predict_classes(X_test)
2 #Accuracy with the test data
3 from sklearn.metrics import accuracy_score
4 print(accuracy_score(labels, pred_cnn_pr2))
```

`model.predict_classes()` is deprecated and will be removed after 2021-01-01. Please use instead:
* `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation).
* `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).

0.9277117973079968

In [106]:

```
1 acc = history_ag.history['accuracy']
2 val_acc = history_ag.history['val_accuracy']
3
4 loss = history_ag.history['loss']
5 val_loss = history_ag.history['val_loss']
6
7 epochs_range = range(15)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```



In [107]:

```
1 val_performance_pr['CNN_1']= score_cnn_pr2[1]
2 performance_pr['CNN_1']=accuracy_score(labels, pred_cnn_pr2)
3 Model_pr['CNN_1']='CNN_1'
```

CNN_2

```
In [108]: 1 img_height = 60
           2 img_width = 60
```

```
In [109]: 1 cnn_i_model = tf.keras.Sequential(
           2 [
           3     keras.layers.Conv2D(filters=32,kernel_size=(4,4),activation='relu',input_
           4         keras.layers.MaxPool2D(pool_size=(2,2)),# down sampling the output inste_
           5             keras.layers.Conv2D(64, kernel_size=(4, 4), activation='relu'),
           6             keras.layers.Conv2D(64, kernel_size=(4, 4), activation='relu'),
           7             keras.layers.MaxPool2D(pool_size=(2, 2)),
           8             keras.layers.Dropout(0.3),
           9             keras.layers.Flatten(), # flatten out the layers
          10            keras.layers.Dense(256,activation='relu'),
          11            keras.layers.Dense(43,activation = 'softmax')
          12
          13
          14
          15        ])
          16 cnn_i_model.compile(loss='categorical_crossentropy', optimizer='adam', metri
```

```
In [110]: 1 X_train_processed.shape[1:]
```

```
Out[110]: (30, 30, 3)
```

```
In [111]: 1 y_train.shape
```

```
Out[111]: (17625,)
```

```
In [112]: 1 X_val_processed.shape
```

```
Out[112]: (4407, 30, 30, 3)
```

```
In [113]: 1 y_val_cnn.shape
```

```
Out[113]: (4407, 43)
```

```
In [114]: 1 y_train_eq.shape
```

```
Out[114]: (70348,)
```

```
In [115]: 1 y_train_prcnn = to_categorical(y_train_eq, 43)
           2 y_val_prcnn = to_categorical(y_val, 43)
```

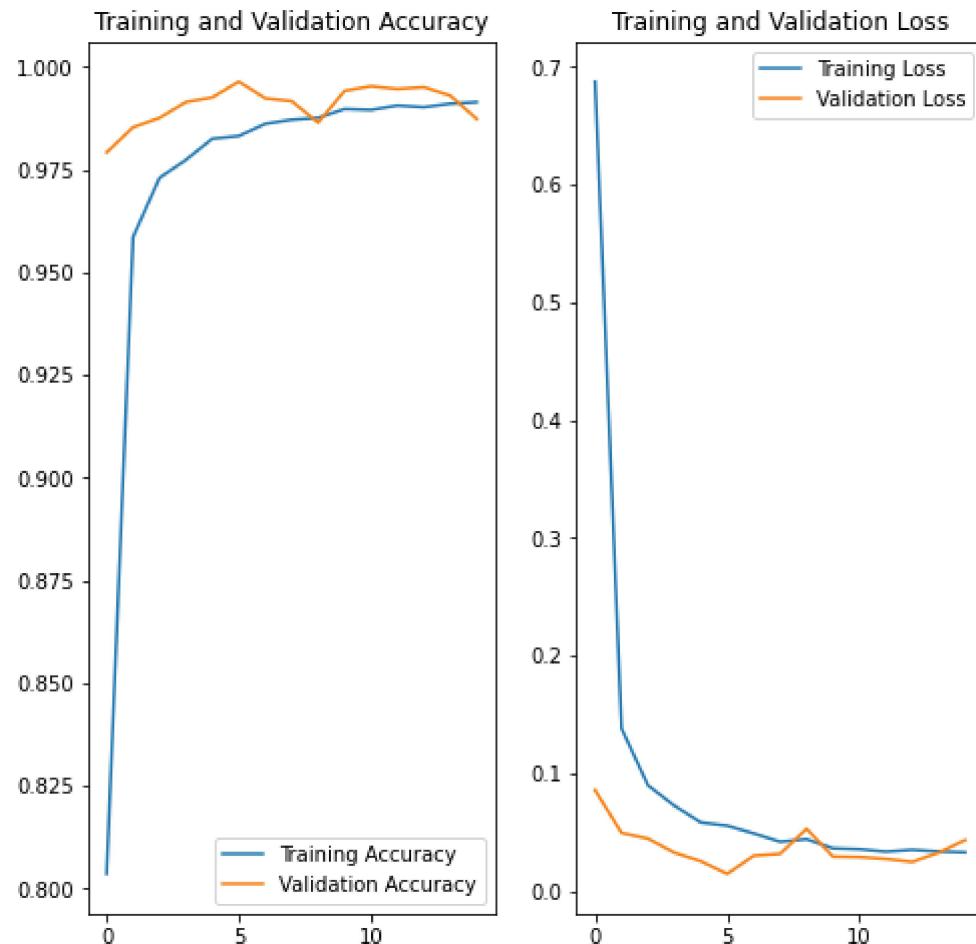
```
In [116]: 1 historya = cnn_i_model.fit(X_train_processed, y_train_prcnn, epochs= 15, va
Epoch 1/15
2199/2199 [=====] - 76s 34ms/step - loss: 1.4670 - accuracy: 0.5919 - val_loss: 0.0858 - val_accuracy: 0.9791
Epoch 2/15
2199/2199 [=====] - 76s 35ms/step - loss: 0.1490 - accuracy: 0.9547 - val_loss: 0.0495 - val_accuracy: 0.9853
Epoch 3/15
2199/2199 [=====] - 76s 34ms/step - loss: 0.0951 - accuracy: 0.9719 - val_loss: 0.0447 - val_accuracy: 0.9875
Epoch 4/15
2199/2199 [=====] - 77s 35ms/step - loss: 0.0748 - accuracy: 0.9765 - val_loss: 0.0327 - val_accuracy: 0.9914
Epoch 5/15
2199/2199 [=====] - 76s 34ms/step - loss: 0.0568 - accuracy: 0.9832 - val_loss: 0.0255 - val_accuracy: 0.9925
Epoch 6/15
2199/2199 [=====] - 77s 35ms/step - loss: 0.0569 - accuracy: 0.9826 - val_loss: 0.0147 - val_accuracy: 0.9964
Epoch 7/15
2199/2199 [=====] - 77s 35ms/step - loss: 0.0417 - accuracy: 0.9879 - val_loss: 0.0301 - val_accuracy: 0.9923
Epoch 8/15
2199/2199 [=====] - 76s 34ms/step - loss: 0.0391 - accuracy: 0.9878 - val_loss: 0.0318 - val_accuracy: 0.9916
Epoch 9/15
2199/2199 [=====] - 76s 35ms/step - loss: 0.0393 - accuracy: 0.9886 - val_loss: 0.0530 - val_accuracy: 0.9864
Epoch 10/15
2199/2199 [=====] - 78s 36ms/step - loss: 0.0363 - accuracy: 0.9897 - val_loss: 0.0294 - val_accuracy: 0.9941
Epoch 11/15
2199/2199 [=====] - 79s 36ms/step - loss: 0.0298 - accuracy: 0.9911 - val_loss: 0.0290 - val_accuracy: 0.9952
Epoch 12/15
2199/2199 [=====] - 77s 35ms/step - loss: 0.0305 - accuracy: 0.9908 - val_loss: 0.0274 - val_accuracy: 0.9946
Epoch 13/15
2199/2199 [=====] - 77s 35ms/step - loss: 0.0314 - accuracy: 0.9903 - val_loss: 0.0249 - val_accuracy: 0.9950
Epoch 14/15
2199/2199 [=====] - 76s 35ms/step - loss: 0.0315 - accuracy: 0.9917 - val_loss: 0.0328 - val_accuracy: 0.9930
Epoch 15/15
2199/2199 [=====] - 76s 35ms/step - loss: 0.0333 - accuracy: 0.9908 - val_loss: 0.0435 - val_accuracy: 0.9873
```

```
In [117]: 1 score_cnn_pr1 = cnn_i_model.evaluate(X_val, y_val_cnn,batch_size=1, verbose=
2 print('Test loss:', score_cnn_pr1[0], ' Test accuracy:', score_cnn_pr1[1])
```

Test loss: 30.481815338134766 Test accuracy: 0.9260267615318298

In [118]:

```
1 acc = historya.history['accuracy']
2 val_acc = historya.history['val_accuracy']
3
4 loss = historya.history['loss']
5 val_loss = historya.history['val_loss']
6
7 epochs_range = range(15)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```



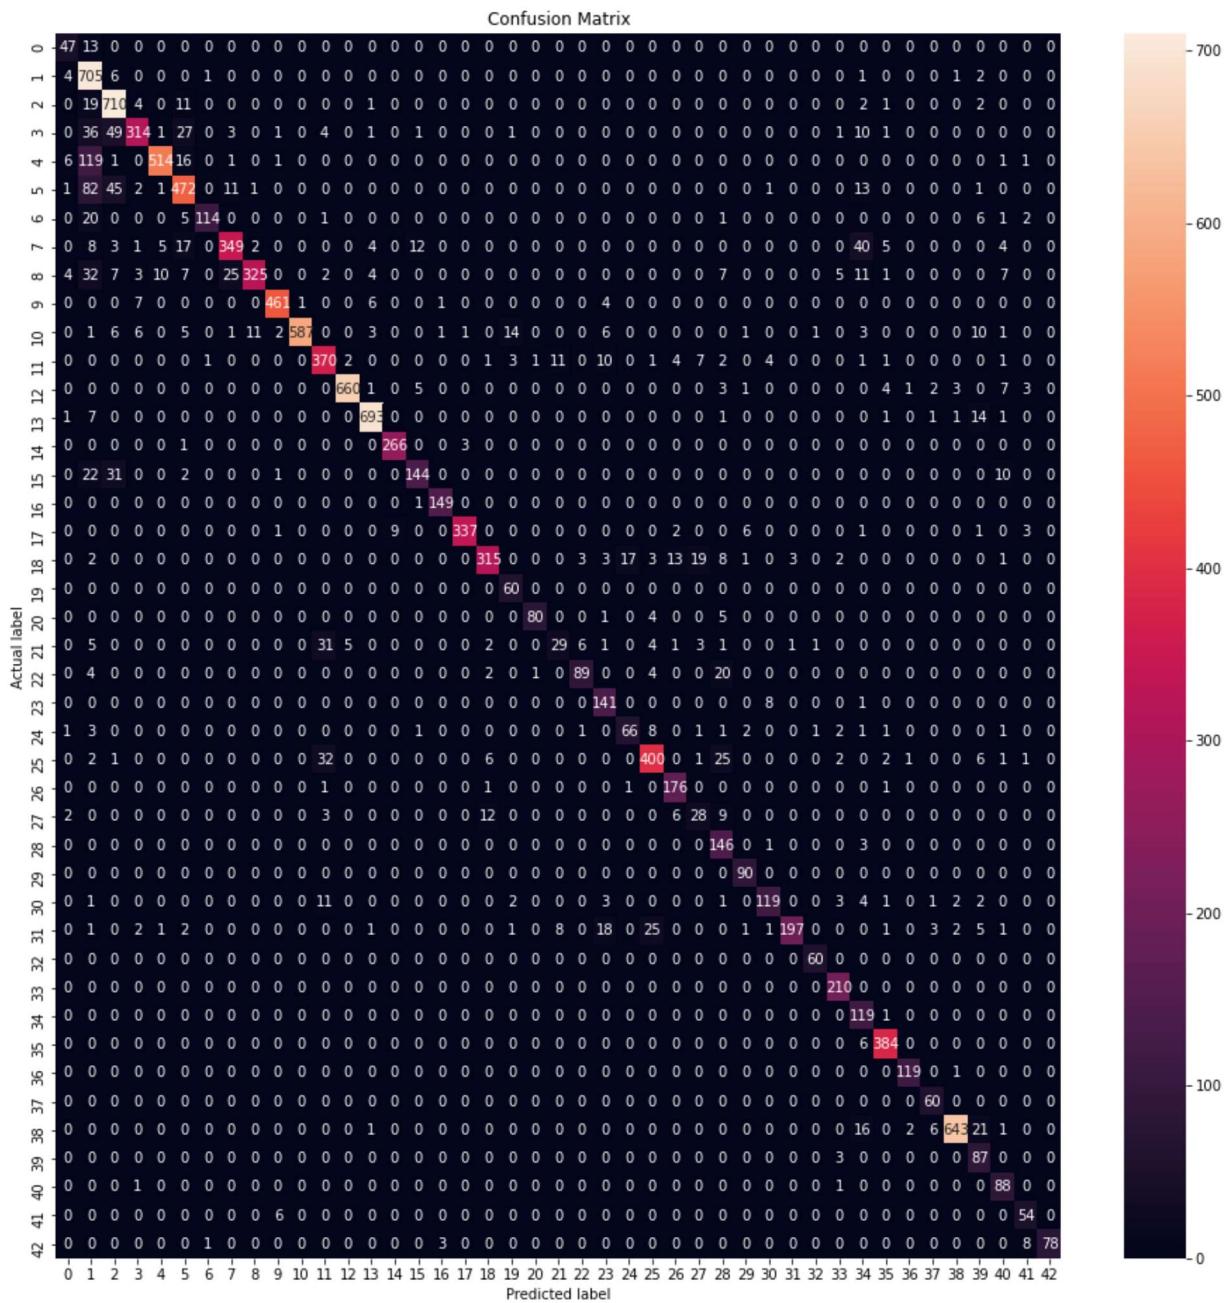
In [119]:

```
1 pred_cnn_pr1 = cnn_i_model.predict_classes(X_test)
2 #Accuracy with the test data
3 from sklearn.metrics import accuracy_score
4 print(accuracy_score(labels, pred_cnn_pr1))
```

0.8752969121140143

In [121]:

```
1 import sklearn
2 import seaborn as sns
3 from sklearn.metrics import confusion_matrix
4 def plot_cm(labels, pred):
5     cm = confusion_matrix(labels, pred)
6     plt.figure(figsize=(16,16))
7     sns.heatmap(cm, annot=True, fmt="d")
8     plt.title('Confusion Matrix')
9     plt.ylabel('Actual label')
10    plt.xlabel('Predicted label')
11
12 plot_cm(labels, pred_cnn_pr1)
```



In [122]:

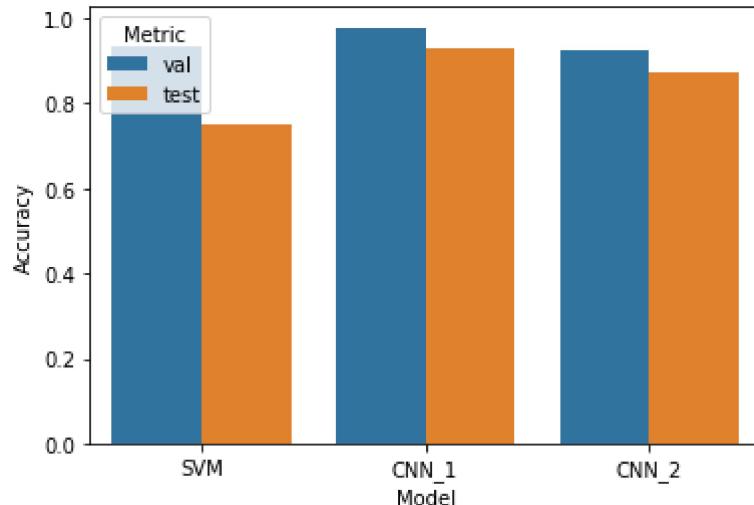
```
1 val_performance_pr['CNN_2']= score_cnn_pr1[1]
2 performance_pr['CNN_2']=accuracy_score(labels, pred_cnn_pr1)
3 Model_pr['CNN_2']='CNN_2'
```

Model Performance

```
In [123]: 1 Performance_pr = pd.DataFrame(  
2     {'val': val_performance_pr,  
3      'test': performance_pr,  
4      'Model':Model_pr  
5 })
```

```
In [124]: 1 df_accuracy_pr = Performance_pr.melt(id_vars=['Model'], var_name='Metric', v
```

```
In [125]: 1 ax = sns.barplot(x="Model", y="Accuracy", hue="Metric", data=df_accuracy_pr)
```



Explainable AI

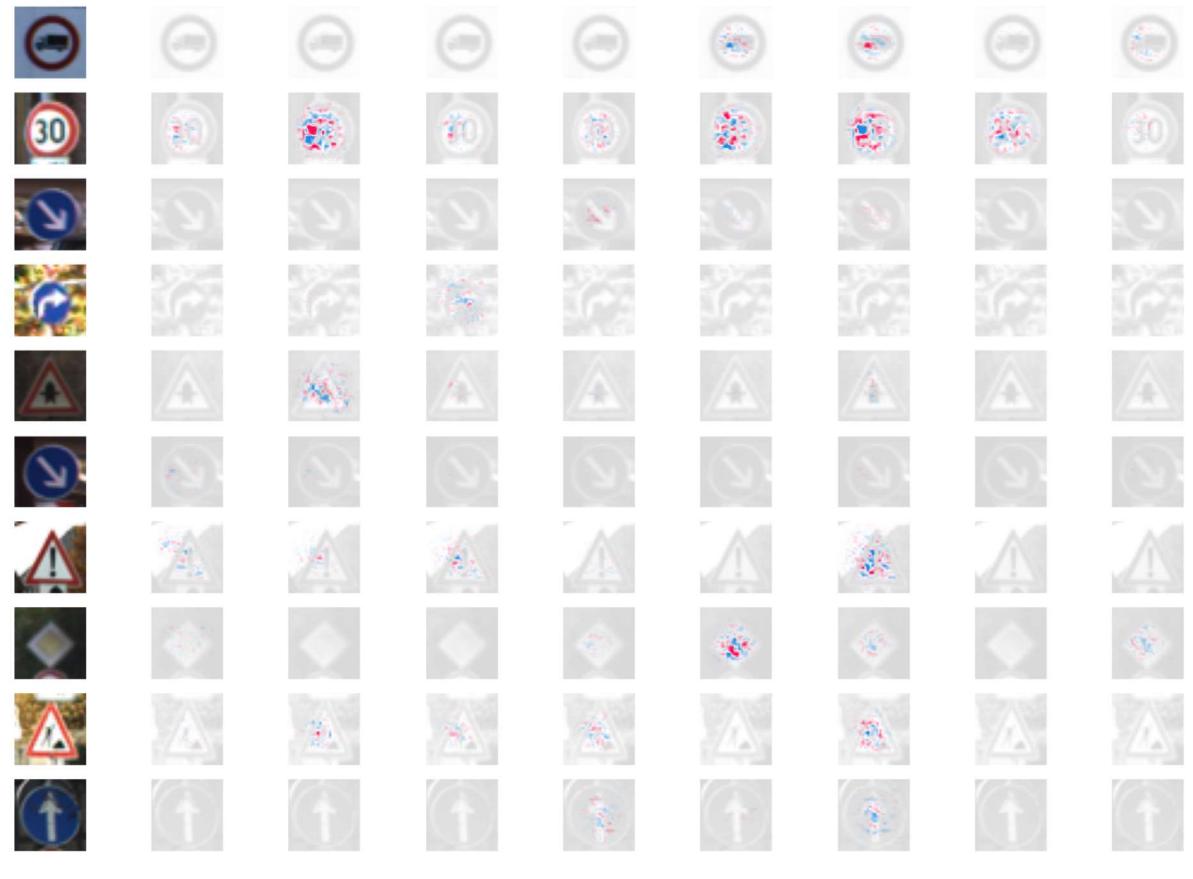
```
In [181]: 1 from sklearn.metrics import accuracy_score  
2 y_test_at = pd.read_csv("C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Anal  
3 labels = y_test_at["ClassId"].values  
4 imgs = y_test_at["Path"].values  
5 data=[]  
6 for img in imgs:  
7     image = Image.open("C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Analy  
8     image = image.resize((30,30))  
9     image = np.array(image)  
10    image = image.astype('float32')  
11    image /=255.0  
12    data.append(np.array(image))  
13 X_test_at=np.array(data)
```

```
In [182]: 1 explainer = shap.GradientExplainer(cnn_ag_model, X_train_processed)
```

```
In [183]: 1 # we explain the model's predictions on the first three samples of the test  
2 shap_values = explainer.shap_values(X_test_at[:10])
```

`tf.keras.backend.set_learning_phase` is deprecated and will be removed after 2020-10-11. To update it, simply pass a True/False value to the `training` argument of the `__call__` method of your layer or model.

```
In [185]: 1 # here we plot the explanations for all classes for the first input (this is  
2 shap.image_plot([shap_values[i] for i in range(8)], X_test_at[:10])
```



Object Detection

```
In [142]: 1 from tqdm import tqdm
2 import math
3 #Run this only once
4 #Creating validation directory
5 val_dir = 'C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Analytics\\\\Final P
6 os.mkdir(val_dir)
7
8 n_classes = 43
9 train_dir = 'C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Analytics\\\\Final
10
11 #Moving files from train to validation directory
12 for n in tqdm(range(n_classes)) :
13     path = os.path.join(val_dir, str(n))
14     os.mkdir(path)
15     src_path = train_dir.format('/' + str(n))
16     files = os.listdir(src_path)
17     rand_idx = random.sample(range(len(files)), math.ceil(len(files)/4))
18     for idx in rand_idx :
19         src = src_path + "/" + files[idx]
20         shutil.move(src, path)
```

100% | 43/43 [00:05<00:00, 7.40it/s]

```
In [143]: 1 #Setting up variables
2 IMG_WIDTH = 30
3 IMG_HEIGHT = 30
4 N_CHANNELS = 3
5 BATCH_SIZE = 32
6 N_EPOCHS = 10
7 VAL_BATCH_SIZE = 32
8 CLASS_NAMES = list(range(43))
9 N_CLASSES = 43
10
11 train_path = "C:/Drive/Siri/MSDSBA/4. IE7860 Intelligent Analytics/Final Pro
12 val_path = "C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Analytics\\\\Final
```

In [144]:

```
1 import pathlib
2 import os
3 import ntpath
4 #Path to train and validation datasets
5 data_root_train = pathlib.Path(train_path)
6 data_root_val = pathlib.Path(val_path)
7
8 #Getting paths to all the images in train and validation sets
9 all_image_paths_train = list(data_root_train.glob('*/*'))
10 all_image_paths_train = [str(path).replace(os.sep,'/')] for path in all_image_
11
12 all_image_paths_val = list(data_root_val.glob('*/*'))
13 all_image_paths_val = [str(path).replace(os.sep,'/')] for path in all_image_p
14
15 #Counting number of images in each sets
16 image_count_train = len(all_image_paths_train)
17 image_count_val = len(all_image_paths_val)
18 all_image_paths_train
```

Out[144]:

In [145]:

```

1 #Extracting labels for each image
2 label_names_train = sorted(int(item.name) for item in data_root_train.glob('*'))
3 label_names_val = sorted(int(item.name) for item in data_root_val.glob('*'))
4 label_to_index_train = dict((name, index) for index, name in enumerate(label_names_train))
5 label_to_index_val = dict((name, index) for index, name in enumerate(label_names_val))
6 all_image_labels_train = [label_to_index_train[int(pathlib.Path(path).parent.name)] for path in data_root_train.glob('*')]
7 all_image_labels_val = [label_to_index_val[int(pathlib.Path(path).parent.name)] for path in data_root_val.glob('*')]

```

```
In [146]: 1 df_train = pd.read_csv("C:\\\\Drive\\\\Siri\\\\MSDSBA\\\\4. IE7860 Intelligent Analy
2 for idx, row in df_train.iterrows() :
3     w = row['Width']
4     h = row['Height']
5     if w > IMG_WIDTH :
6         diff = w-IMG_WIDTH
7         df_train.iloc[idx, 4] = df_train.iloc[idx]['Roi.X2'] - diff
8     else :
9         diff = IMG_WIDTH-w
10        df_train.iloc[idx, 4] = df_train.iloc[idx]['Roi.X2'] + diff
11    if h > IMG_HEIGHT :
12        diff = h - IMG_HEIGHT
13        df_train.iloc[idx, 5] = df_train.iloc[idx]['Roi.Y2'] - diff
14    else :
15        diff = IMG_HEIGHT - h
16        df_train.iloc[idx, 5] = df_train.iloc[idx]['Roi.Y2'] + diff
```

```
In [147]: 1 df_train
```

Out[147]:

	Width	Height	Roi.X1	Roi.Y1	Roi.X2	Roi.Y2	ClassId	Path
0	27	26	5	5	25	24	20	Train/20/00020_00000_00000.png
1	28	27	5	6	25	25	20	Train/20/00020_00000_00001.png
2	29	26	6	5	25	25	20	Train/20/00020_00000_00002.png
3	28	27	5	6	25	25	20	Train/20/00020_00000_00003.png
4	28	26	5	5	25	25	20	Train/20/00020_00000_00004.png
...
39204	52	56	5	6	25	25	42	Train/42/00042_00007_00025.png
39205	56	58	5	5	25	25	42	Train/42/00042_00007_00026.png
39206	58	62	5	6	25	25	42	Train/42/00042_00007_00027.png
39207	63	69	5	7	25	24	42	Train/42/00042_00007_00028.png
39208	68	69	7	6	24	24	42	Train/42/00042_00007_00029.png

39209 rows × 8 columns

```
In [148]: 1 all_image_paths_val[6][83:]
```

Out[148]: '0/00000_00002_00014.png'

```
In [149]:  
1 train_idx_list = []  
2 val_idx_list = []  
3  
4 for path_tr in tqdm(all_image_paths_train) :  
5     train_idx_list.append(df_train[df_train['Path'] == path_tr[72 : ]].index[0])  
6 for path_val in tqdm(all_image_paths_val) :  
7     path_val = "Train/" + path_val[83:]  
8     val_idx_list.append(df_train[df_train['Path'] == path_val].index[0])  
9 new_df_train = pd.DataFrame()  
10 new_df_val = pd.DataFrame()  
11 new_df_train = new_df_train.append(df_train.iloc[train_idx_list], ignore_index=True)  
12 new_df_val = new_df_val.append(df_train.iloc[val_idx_list], ignore_index=True)  
13 new_df_train = new_df_train.drop(['Height', 'Width', 'ClassId', 'Path'], axis=1)  
14 new_df_val = new_df_val.drop(['Height', 'Width', 'ClassId', 'Path'], axis=1)
```

100%|██████████| 16509/16509 [00:51<00:00, 319.54it/s]
100%|██████████| 5523/5523 [00:18<00:00, 299.23it/s]

```
In [150]: 1 new_df_val
```

Out[150]:

	Roi.X1	Roi.Y1	Roi.X2	Roi.Y2
0	5	6	25	24
1	6	6	25	25
2	6	5	25	25
3	10	11	19	19
4	12	13	18	18
...
5518	7	7	24	24
5519	6	6	25	24
5520	5	6	25	25
5521	5	6	25	25
5522	5	5	24	24

5523 rows × 4 columns

```

In [151]: 1 def tfdata_generator(images, labels, df, is_training, batch_size=32):
2     '''Construct a data generator using tf.Dataset'''
3     def parse_function(filename, labels, df):
4         '''Function to preprocess the images'''
5         #reading path
6         image_string = tf.io.read_file(filename)
7         #decoding image
8         image = tf.image.decode_png(image_string, channels=N_CHANNELS)
9         # This will convert to float values in [0, 1]
10        image = tf.image.convert_image_dtype(image, tf.float32)
11        #Adjusting contrast and brightness of the image
12        if tf.math.reduce_mean(image) < 0.3 :
13            image = tf.image.adjust_contrast(image, 5)
14            image = tf.image.adjust_brightness(image, 0.2)
15        #resize the image
16        image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH], method="nearest")
17        image = image/255.0
18        #one hot coding for Label
19        #y = tf.one_hot(tf.cast(label, tf.uint8), N_CLASSES)
20        return image, {"classification" : labels, "regression" : df}
21    ##creating a dataset from tensorslices
22    dataset = tf.data.Dataset.from_tensor_slices((images, labels, df))
23    if is_training:
24        dataset = dataset.shuffle(30000) # depends on sample size
25        # Transform and batch data at the same time
26        dataset = dataset.map(parse_function, num_parallel_calls = tf.data.experimental.AUTOTUNE)
27        dataset = dataset.repeat()
28        dataset = dataset.batch(batch_size)
29        #prefetch the data into CPU/GPU
30        dataset = dataset.prefetch(tf.data.experimental.AUTOTUNE)
31    return dataset

```

```

In [152]: 1 tf_image_generator_train = tfdata_generator(all_image_paths_train, all_image_labels_train)
2 tf_image_generator_val = tfdata_generator(all_image_paths_val, all_image_labels_val)

```

```

In [153]: 1 steps_per_epoch_train = np.ceil(len(all_image_paths_train)/BATCH_SIZE)
2 steps_per_epoch_val = np.ceil(len(all_image_paths_val)/BATCH_SIZE)

```

```

In [154]: 1 class Sharpen(tf.keras.layers.Layer):
2     def __init__(self, num_outputs) :
3         super(Sharpen, self).__init__()
4         self.num_outputs = num_outputs
5
6         def build(self, input_shape) :
7             self.kernel = np.array([[-1,-1,-1], [-1,9,-1], [-1,-1,-1]])
8             self.kernel = tf.expand_dims(self.kernel, 0)
9             self.kernel = tf.expand_dims(self.kernel, 0)
10            self.kernel = tf.cast(self.kernel, tf.float32)
11
12            def call(self, input_) :
13                return tf.nn.conv2d(input_, self.kernel, strides=[1, 1, 1, 1], padding='VALID')

```

In [161]:

```
1 def get_model() :
2     #Input layer
3     input_layer = Input(shape=(IMG_HEIGHT, IMG_WIDTH, N_CHANNELS, ), name="inp")
4     #Sharpen Layer to sharpen the edges of the image.
5     sharp = Sharpen(num_outputs=(IMG_HEIGHT, IMG_WIDTH, N_CHANNELS, ))(input_l)
6     #Convolution, maxpool and dropout layers
7     conv_1 = Conv2D(filters=32, kernel_size=(5,5), activation=relu,
8                     kernel_initializer=he_normal(seed=54), bias_initializer=ze
9                     name="first_convolutional_layer") (sharp)
10    conv_2 = Conv2D(filters=64, kernel_size=(3,3), activation=relu,
11                     kernel_initializer=he_normal(seed=55), bias_initializer=ze
12                     name="second_convolutional_layer") (conv_1)
13    maxpool_1 = MaxPool2D(pool_size=(2,2), name = "first_maxpool_layer")(conv_
14    dr1 = Dropout(0.25)(maxpool_1)
15    conv_3 = Conv2D(filters=64, kernel_size=(3,3), activation=relu,
16                     kernel_initializer=he_normal(seed=56), bias_initializer=ze
17                     name="third_convolutional_layer") (dr1)
18    maxpool_2 = MaxPool2D(pool_size=(2,2), name = "second_maxpool_layer")(conv_
19    dr2 = Dropout(0.25)(maxpool_2)
20    flat = Flatten(name="flatten_layer")(dr2)
21
22     #Fully connected Layers
23    d1 = Dense(units=256, activation=relu, kernel_initializer=he_normal(seed=4
24                     bias_initializer=zeros(), name="first_dense_layer_classificatio
25    dr3 = Dropout(0.5)(d1)
26
27    classification = Dense(units = 43, activation=None, name="classification",
28
29    regression = Dense(units = 4, activation = 'linear', name = "regression",
30                     kernel_initializer=RandomNormal(seed=43), kernel_regula
31     #Model
32    model = Model(inputs = input_layer, outputs = [classification, regression]
33    model.summary()
34    return model
```

```
In [166]: 1 model = get_model()
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	[None, 30, 30, 3]	0	
sharpen_3 (Sharpen)	(None, 30, 30, 3)	0	input_layer[0]
first_convolutional_layer (Conv)	(None, 26, 26, 32)	2432	sharpen_3[0]
second_convolutional_layer (Con)	(None, 24, 24, 64)	18496	first_convolutional_layer[0][0]
first_maxpool_layer (MaxPooling)	(None, 12, 12, 64)	0	second_convolutional_layer[0][0]
dropout_7 (Dropout)	(None, 12, 12, 64)	0	first_maxpool_layer[0][0]
third_convolutional_layer (Conv)	(None, 10, 10, 64)	36928	dropout_7[0]
second_maxpool_layer (MaxPoolin)	(None, 5, 5, 64)	0	third_convolutional_layer[0][0]
dropout_8 (Dropout)	(None, 5, 5, 64)	0	second_maxpool_layer[0][0]
flatten_layer (Flatten)	(None, 1600)	0	dropout_8[0]
first_dense_layer_classification	(None, 256)	409856	flatten_layer[0][0]
dropout_9 (Dropout)	(None, 256)	0	first_dense_layer_classification[
classification (Dense)	(None, 43)	11051	dropout_9[0]

```
regression (Dense)           (None, 4)      1028      dropout_9[0]
[0]
=====
=====
Total params: 479,791
Trainable params: 479,791
Non-trainable params: 0
```

In [167]:

```
1  from tensorflow.keras.callbacks import Callback
2  from sklearn.metrics import f1_score, precision_score, recall_score
3  class Metrics(Callback) :
4      def __init__(self, validation_data_generator) :
5          self.validation_data_generator = validation_data_generator
6
7      def on_train_begin(self, logs={}) :
8          '''
9              This function initializes lists to store AUC and Micro F1 scores
10             '''
11         self.val_f1s = []
12         self.val_precisions = []
13         self.val_recalls = []
14         self.batches = self.validation_data_generator.as_numpy_iterator()
15
16      def on_epoch_end(self, epoch, logs = {}) :
17          '''
18              This function calculates the micro f1 and auc scores
19              at the end of each epochs
20              '''
21         current_batch = self.batches.next()
22         images = current_batch[0]
23         labels = current_batch[1]
24         labels = labels["classification"]
25         labels = np.array(labels)
26         pred = self.model.predict(images)
27         pred = pred[0]
28         val_predict = (np.asarray(pred)).round()
29         idx = np.argmax(val_predict, axis=-1)
30         a = np.zeros( val_predict.shape )
31         a[ np.arange(a.shape[0]), idx ] = 1
32         val_predict = [np.where(r==1)[0][0] for r in a]
33         val_predict = np.array(val_predict)
34         val_targ = labels
35         _val_f1 = f1_score(val_targ, val_predict, average = 'weighted')
36         _val_precision = precision_score(val_targ, val_predict, average='weighted')
37         _val_recall = recall_score(val_targ, val_predict, average='weighted')
38         print("\nEpoch : {0} - Precision_Score : {1:.2f} - Recall_Score : {2:.2f}"
39         self.val_f1s.append(_val_f1)
40         self.val_precisions.append(_val_precision)
41         self.val_recalls.append(_val_recall)
42         return
```

In [168]:

```
1 from keras import backend as K
2 def r2_keras(y_true, y_pred):
3     SS_res = K.sum(K.square(y_true - y_pred))
4     SS_tot = K.sum(K.square(y_true - K.mean(y_true)))
5     return ( 1 - SS_res/(SS_tot + K.epsilon()) )
6
7 loss = SparseCategoricalCrossentropy(from_logits=True)
8
9 model.compile(optimizer="adam", loss = {"classification" : loss, "regression"
10                         metrics={"classification" : "acc", "regression" : r2_keras}, 1
```

In [169]:

```
1 from tensorflow.keras.callbacks import ModelCheckpoint, TensorBoard, EarlySt
2 from sklearn.metrics import f1_score, precision_score, recall_score, auc
3 import datetime
4 import time
5 %load_ext tensorboard
6 log_dir="/content/drive/My Drive/CaseStudy2/logs/fit/" + datetime.datetime.n
7 NAME = "TrafficSignRecog-first-cut-{0}".format(int(time.time()))
8 save_best_model = ModelCheckpoint(filepath='/content/drive/My Drive/CaseStud
9                                         save_best_only = True, mode = 'min', save_
10 tensorboard_callback = TensorBoard(log_dir=log_dir, histogram_freq=1, write_
11                                         write_images = True)
12 early_stop = EarlyStopping(monitor='loss', min_delta=0.0001, patience=5)
13 metrics = Metrics(tf_image_generator_val)
```

In [170]:

```
1 history = model.fit_generator(
2     generator = tf_image_generator_train, steps_per_epoch = steps_per_epoch_
3     epochs = N_EPOCHS,
4     validation_data = tf_image_generator_val, validation_steps = steps_per_e
5     callbacks = [save_best_model, tensorboard_callback, metrics, early_stop]
6 )
```

WARNING:tensorflow:Model failed to serialize as JSON. Ignoring... Layer Sharp
en has arguments in `__init__` and therefore must override `get_config`.

`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.

```
In [171]: 1 test_df = pd.read_csv('C:/Drive/Siri/MSDSBA/4. IE7860 Intelligent Analytics/2018-09-18_10-45-00.csv')
2 for idx, row in test_df.iterrows() :
3     w = row['Width']
4     h = row['Height']
5     if w > IMG_WIDTH :
6         diff = w-IMG_WIDTH
7         test_df.iloc[idx, 4] = test_df.iloc[idx]['Roi.X2'] - diff
8     else :
9         diff = IMG_WIDTH-w
10        test_df.iloc[idx, 4] = test_df.iloc[idx]['Roi.X2'] + diff
11    if h > IMG_HEIGHT :
12        diff = h - IMG_HEIGHT
13        test_df.iloc[idx, 5] = test_df.iloc[idx]['Roi.Y2'] - diff
14    else :
15        diff = IMG_HEIGHT - h
16        test_df.iloc[idx, 5] = test_df.iloc[idx]['Roi.Y2'] + diff
```

```
In [173]: 1 def evaluate_test_images(path, model) :
2     labels = []
3     bbox = []
4     all_imgs = os.listdir(path)
5     all_imgs.sort()
6     for img in tqdm(all_imgs) :
7         if '.png' in img :
8             image_string = tf.io.read_file(path + '/' + img)
9             image = tf.image.decode_png(image_string, channels=N_CHANNELS)
10            image = tf.image.convert_image_dtype(image, tf.float32)
11            if tf.math.reduce_mean(image) < 0.3 :
12                image = tf.image.adjust_contrast(image, 5)
13                image = tf.image.adjust_brightness(image, 0.2)
14                image = tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH], method="nearest")
15                image = image/255.0
16                image = np.expand_dims(image, axis=0)
17                pred = model.predict(image)
18                labels.append(np.argmax(pred[0][0]))
19                bbox.append(pred[1][0])
20    return labels, bbox
```

```
In [174]: 1 model.compile()
```

```
In [175]: 1 test_path = "C:/Drive/Siri/MSDSBA/4. IE7860 Intelligent Analytics/Final Project"
2 labels, bbox = evaluate_test_images(test_path, model)
```

100% [██████████] | 12630/12630 [09:04<00:00, 23.19it/s]

```
In [176]: 1 from sklearn.metrics import f1_score, mean_squared_error, precision_score, recall_score
2 test_precision = precision_score(test_df['ClassId'], labels, average='weighted')
3 test_recall = recall_score(test_df['ClassId'], labels, average='weighted')
4 test_f1_score = f1_score(test_df['ClassId'], labels, average = 'weighted')
5 test_mse = mean_squared_error(test_df.iloc[:, 2:6], bbox)
```

In [177]:

```
1 print("Test Precision : {0:.2f}".format(test_precision))
2 print("Test Recall : {0:.2f}".format(test_recall))
3 print("Test f1-score : {0:.2f}".format(test_f1_score))
4 print("Test Mean squared error : {0:.2f}".format(test_mse))
```

Test Precision : 0.89
Test Recall : 0.88
Test f1-score : 0.87
Test Mean squared error : 3.20

In [178]:

```
1 pip install pillow
```

Requirement already satisfied: pillow in c:\users\peddi\anaconda3\lib\site-packages (8.0.1)
Note: you may need to restart the kernel to use updated packages.

In [179]:

```
1 #Plotting some correctly classified images :
2 from PIL import Image
3 import matplotlib.pyplot as plt
4 import math
5 test_df = pd.read_csv('C:/Drive/Siri/MSDSBA/4. IE7860 Intelligent Analytics/
6 test_path = "C:/Drive/Siri/MSDSBA/4. IE7860 Intelligent Analytics/Final Proj
7 fig, ax = plt.subplots(5, 5)
8 fig.set_size_inches((10, 10))
9 cnt = 0
10 i = 0
11 j = 0
12 while cnt < 25 :
13     idx = random.randint(0, test_df.shape[0])
14     if test_df.iloc[idx]['ClassId'] == labels[idx] :
15         x1, y1, x2, y2 = test_df.iloc[idx, 2:6]
16         img = cv2.imread(test_path + test_df.iloc[idx]['Path'])
17         height = img.shape[0]
18         width = img.shape[1]
19         px1, py1, px2, py2 = bbox[idx][0], bbox[idx][1], bbox[idx][2], bbox[idx]
20         if width > IMG_WIDTH :
21             diff = width - IMG_WIDTH
22             px2 += diff
23         elif width < IMG_WIDTH :
24             px2 -= diff
25         else :
26             pass
27         if height > IMG_HEIGHT :
28             diff = height - IMG_HEIGHT
29             py2 += diff
30         elif height < IMG_HEIGHT :
31             py2 -= diff
32         else :
33             pass
34         img = cv2.rectangle(img, (x1, y1), (x2, y2), (255,0,0), 1) #Actual bound
35         img = cv2.rectangle(img, (math.ceil(px1), math.ceil(py1)), (math.ceil(px
36         ax[i][j].imshow(img)
37         j+=1
38         if j == 5 :
39             i+=1
40             j = 0
41             cnt+=1
42 plt.show()
```

